



DIAG

Reti di Elaboratori

Corso di Laurea in Ingegneria Informatica
Università degli Studi di Roma “La Sapienza”
Prof.ssa Chiara Petrioli

Parte di queste slide sono state prese dal materiale associato al libro
Computer Networking: A Top Down Approach , 7th edition.
All material copyright 1996-2009
J.F Kurose and K.W. Ross, All Rights Reserved
Thanks also to Antonio Capone, Politecnico di Milano, Giuseppe Bianchi and
Francesco LoPresti, Un. di Roma Tor Vergata

Chapter 2 outline

- 2.1 Principles of app layer protocols
 - clients and servers
 - app requirements
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
 - SMTP, POP3, IMAP
- 2.5 DNS
- 2.9 Content distribution
 - Network Web caching
 - Content distribution networks
 - P2P file sharing

I servizi di IP: esempio il WEB

The screenshot shows a vintage-style Netscape browser window. The title bar reads "My Links - Netscape". The menu bar includes "File", "Edit", "View", "Go", "Communicator", and "Help". The toolbar contains icons for Back, Forward, Reload, Home, Search, Netscape, Print, Security, Shop, and Stop. The location bar shows the URL "http://www.elet.polimi.it/Users/DEI/Sections/Telecom/Antonio.Capone/mylinks.html". Below the toolbar is a horizontal bar with links: Bookmarks, Location, Instant Message, WebMail, Radio, People, Yellow Pages, Download, Calendar, Channels, and RealPlayer. The main content area features a logo of a motorcycle and the text "My Links". On the left side, there are two columns of links:

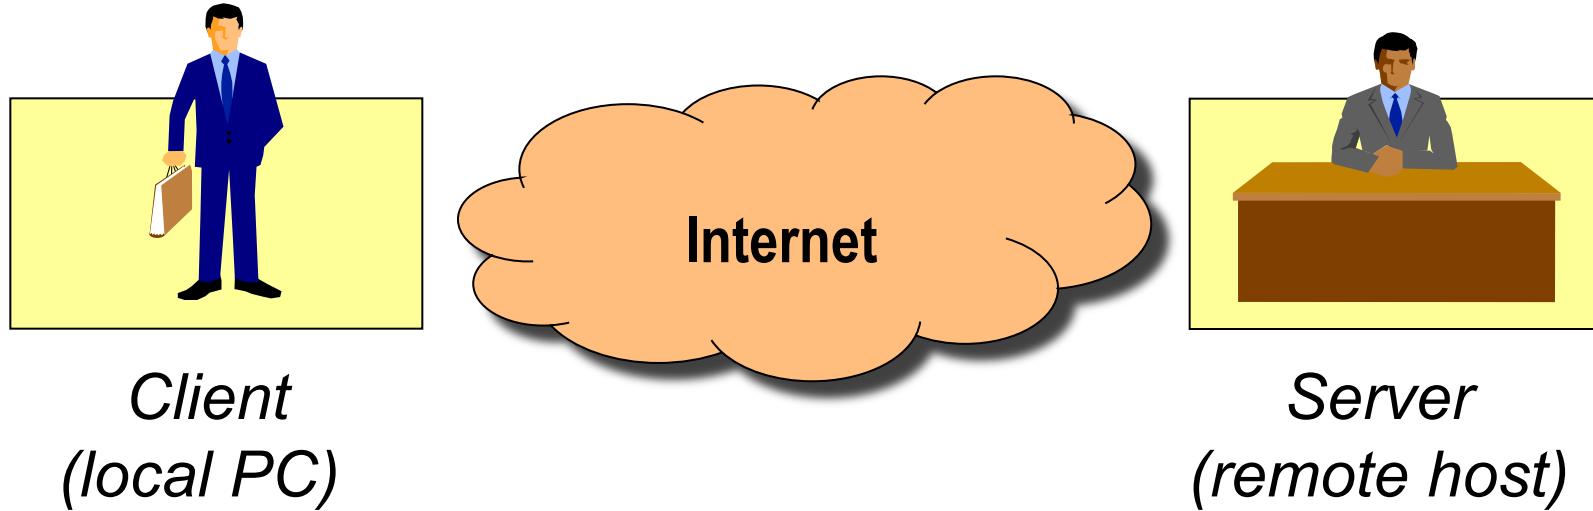
- My Home Page
- My mail box
- Search Engines
- Scientific and Technical Papers search engines
- Conferences, Journals and Technical societies
 - Conferences
 - Journals and Societies
 - On-line magazines
- Research Centers and Organizations
 - Italian Universities
 - Communication Networks
 - Wireless
 - IP

On the right side, there are two columns of links:

- Research Topics
 - Wireless ATM
 - UMTS
 - IP
 - Simulation
- Telecom Companies & info sources
 - Education Centers
 - Telecommunication Networks Documents
 - Software resources
 - Public institutions
 - Art and Culture
 - News Papers
 - Travels
 - Computer Stores
 - Friends Home Pages
 - Weather
 - Italian Internet Access Points

At the bottom, a status bar says "You are offline. Choose 'Go Online...' to connect". The bottom right corner shows the "Application Layer" icon.

Basic scenario



Client wants to retrieve a web page. What happens?

What is a “page” on the web?

a resource (i.e a file), specified by a
URL: Uniform Resource Locator.

e.g. personal web pages:

`HTTP://cerbero.elet.polimi.it/people/bianchi/index.html`

The three components of an URL

1. Protocol (also called "scheme")

- how can a page be accessed? (application protocol used)
 - **http://cerbero.elet.polimi.it/people/bianchi/index.html**

2. Host name

- Where is the page located? (symbolic or numeric location)
 - **http://cerbero.elet.polimi.it/people/bianchi/index.html**

3. File (resource) name

- What is the page called? (with full path)
 - **http://cerbero.elet.polimi.it/people/bianchi/index.html**

Network applications: some jargon

Process: program running within a host.

- within same host, two processes communicate using **interprocess communication** (defined by OS).
- processes running in different hosts communicate with an **application-layer protocol** (defining message format, which message to exchange and in which order)

user agent: interfaces with user “above” and network “below”.

- implements user interface & application-level protocol
 - Web: browser
 - E-mail: mail reader
 - streaming audio/video: media player

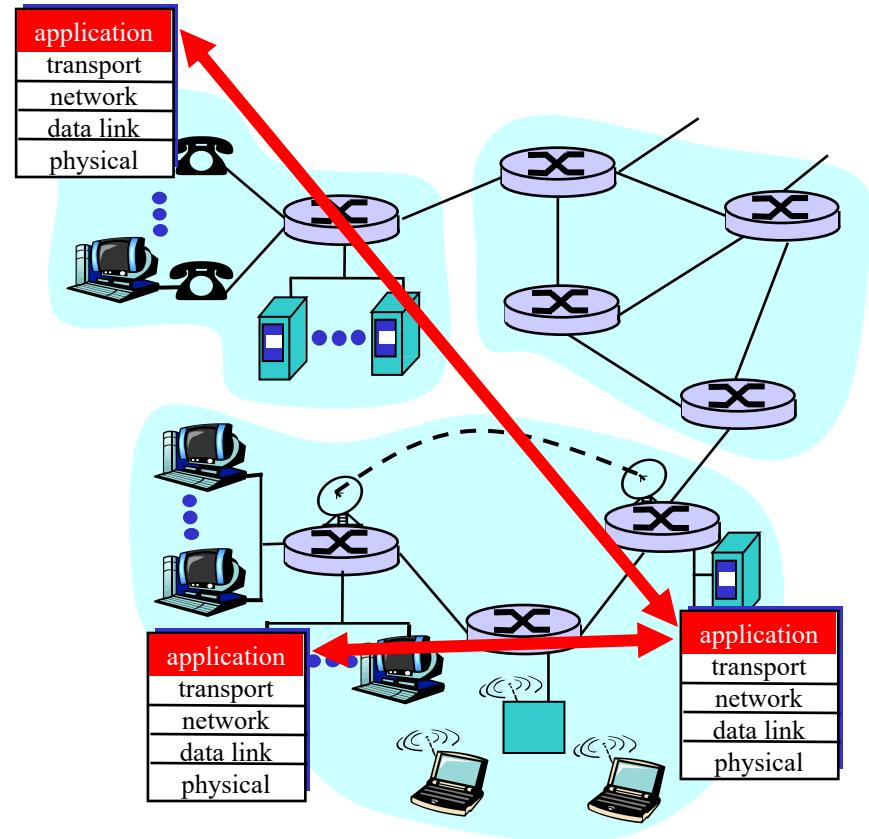
Applications and application-layer protocols

Application: communicating, distributed processes

- e.g., e-mail, Web, P2P file sharing, instant messaging
- running in end systems (hosts)
- exchange messages to implement application

Application-layer protocols

- one “piece” of an app
- define messages exchanged by apps and actions taken
- use communication services provided by lower layer protocols (TCP, UDP)



App-layer protocol defines

- Types of messages exchanged, eg, request & response messages
- Syntax of message types: what fields in messages & how fields are delineated
- Semantics of the fields, ie, meaning of information in fields
- Rules for when and how processes send & respond to messages

Public-domain protocols:

- defined in RFCs
- allows for interoperability
- eg, HTTP, SMTP

Proprietary protocols:

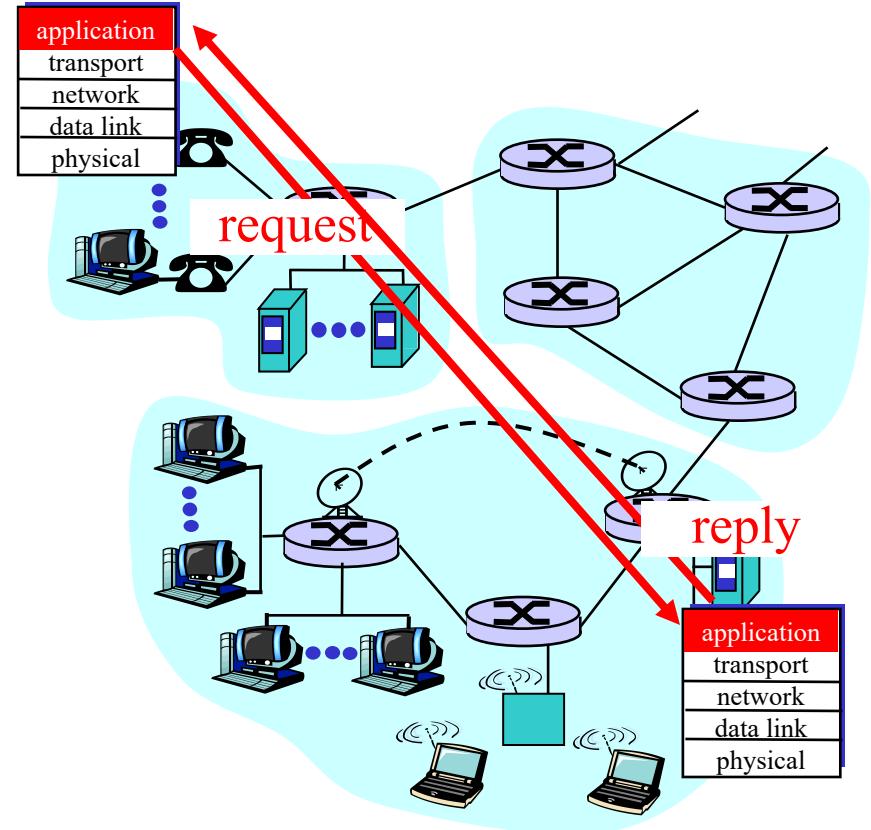
- eg, KaZaA

Client-server paradigm

Typical network app has two pieces:
client and *server*
(NEXT SLIDE)

Client:

- initiates contact with server (“speaks first”)
- typically **requests service** from server,
- Web: client implemented in browser; e-mail: in mail reader

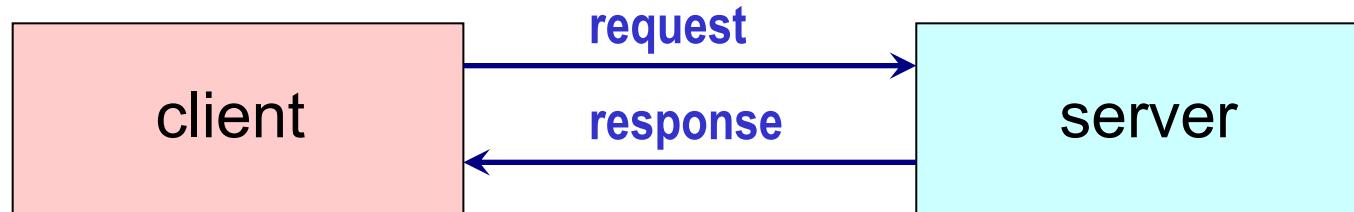


Server:

- provides requested service to client
- e.g., Web server sends requested Web page, mail server delivers e-mail

Architettura Client-Server

- Un processo client è solo in grado fare richieste di servizio (informazioni) e di interpretare le risposte
- Un processo server ha solo il compito di interpretare le richieste e fornire le risposte
- Se è necessario nello stesso host sia fare richieste che fornire risposte vengono usati due processi, client e server.
- Un protocollo applicativo per un'architettura client-server rispecchia questa divisione di ruoli e prevede messaggi di richiesta (request) generati dal lato client e messaggi di risposta (response) generati dal server



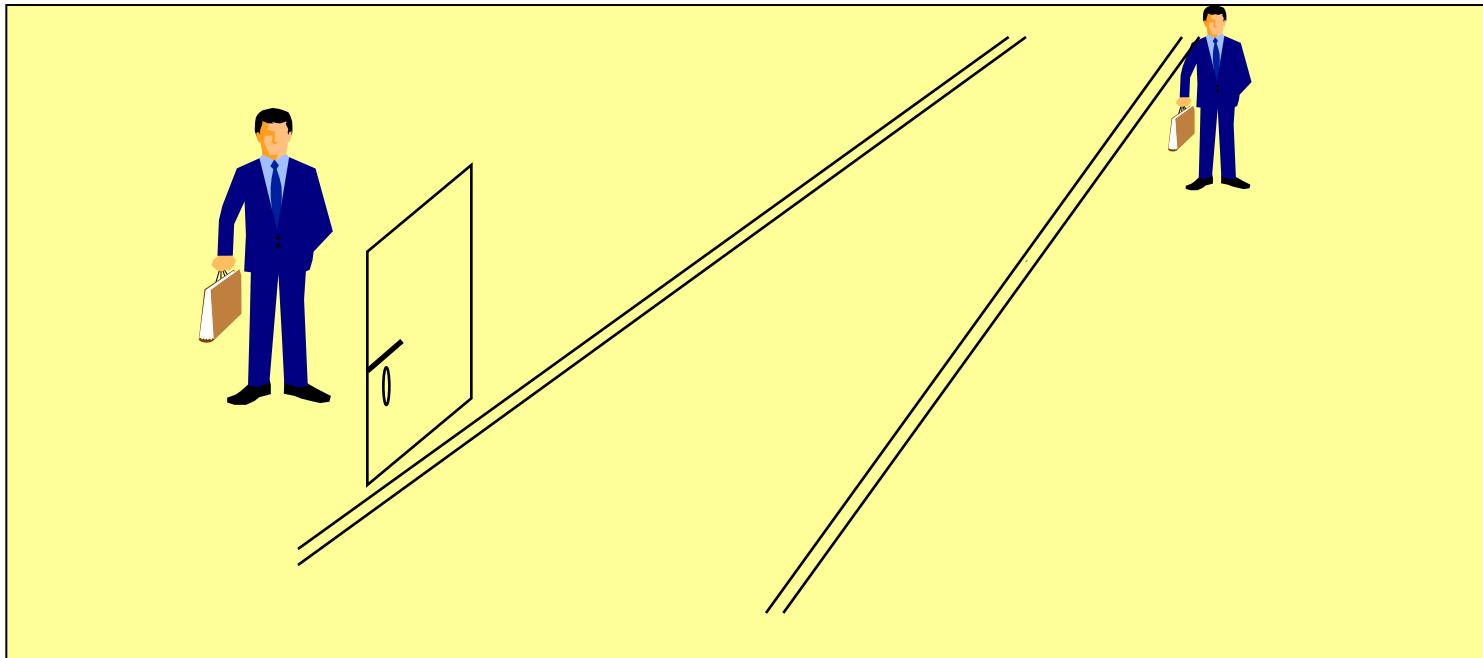
Programmi Client e Server

- Distinguiamo tra programma (software) e processo, istanza del programma in esecuzione su un host
- Un processo server è in esecuzione a tempo illimitato sul proprio host (daemon) e viene attivato mediante una *passive open*
- Un processo client viene attivato solo al momento di fare le richieste e viene attivato mediante una *active open* su richiesta dell'utente o di altro processo applicativo
- la passive open del server fa sì che da quel momento il server accetti richieste dai client
- **Applicazione Web:** l'active open del client richiede l'indicazione dell'indirizzo e della porta del client, ed inizia la connessione TCP per mettere in comunicazione il client e il server

Addressing processes:

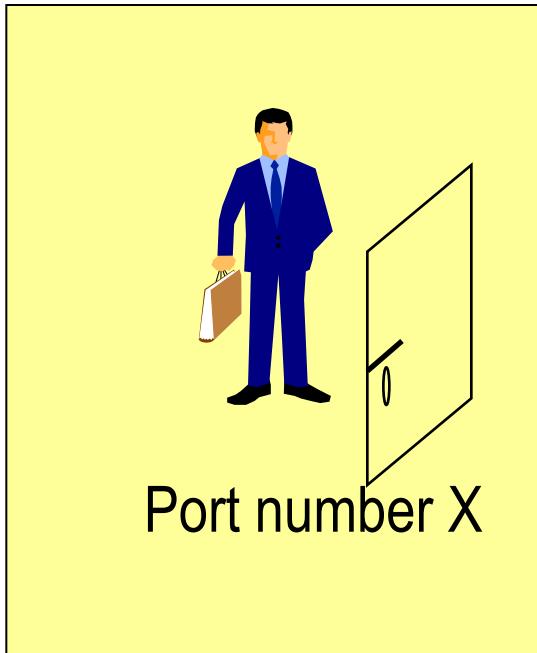
- For a process to receive messages, it must have an identifier
- Every host has a unique 32-bit IP address
- Q: does the IP address of the host on which the process runs suffice for identifying the process?
- Answer: No, many processes can be running on same host
- Identifier includes both the IP address and **port numbers** associated with the process on the host.
- Example port numbers:
 - HTTP server: 80
 - Mail server: 25

Port numbers

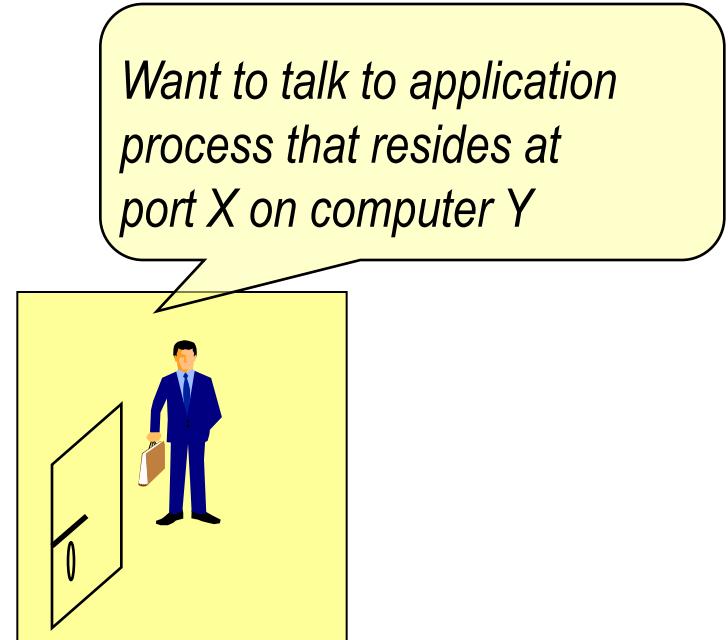


the “address” of the SW process inside the computer!

Same addressing scheme works for different machines!!



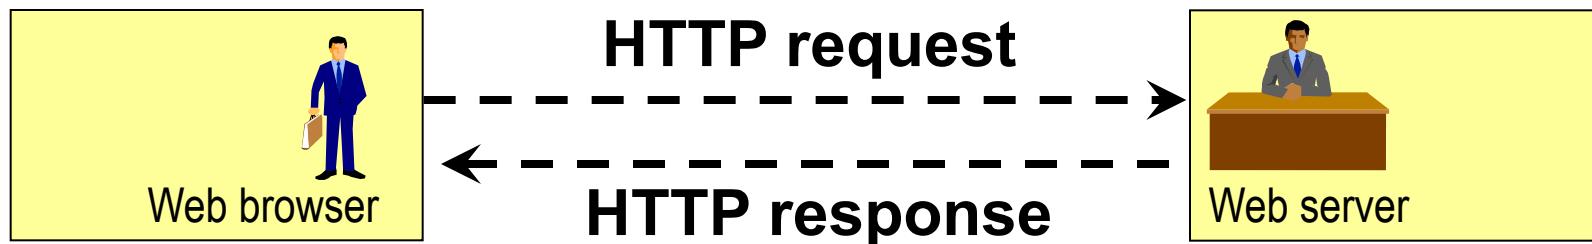
Computer address Y



Re-understanding URLs

`HTTP://cerbero.elet.polimi.it/people/bianchi/index.html`

protocol location Filename



Web browser (SW process) needs to send HTTP Request to Web Server.
Location not enough (it is just the computer address!)

Addressing web servers: (wrong) idea

HTTP://cerbero.elet.polimi.it/....

The diagram shows the URL "HTTP://cerbero.elet.polimi.it/...." with three curly braces underneath it. The first brace covers "HTTP:", labeled "protocol". The second brace covers "cerbero.elet.polimi.it", labeled "location". The third brace covers the trailing slashes and dots, labeled "Filename".

protocol location Filename

Location = computer address!

Protocol = SW process address (HTTP = goto Web Server)

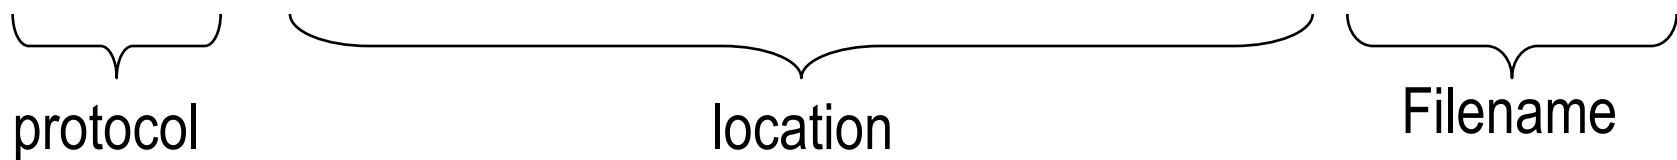
.... **No need for port numbers??**

...

What if more than one Web Server installed???

URL structure (corrected!)

HTTP://cerbero.elet.polimi.it : *portnumber* / ...



Default value for HTTP: 80

HTTP://cerbero.elet.polimi.it/...

equal to

HTTP://cerbero.elet.polimi.it:80/...

different from

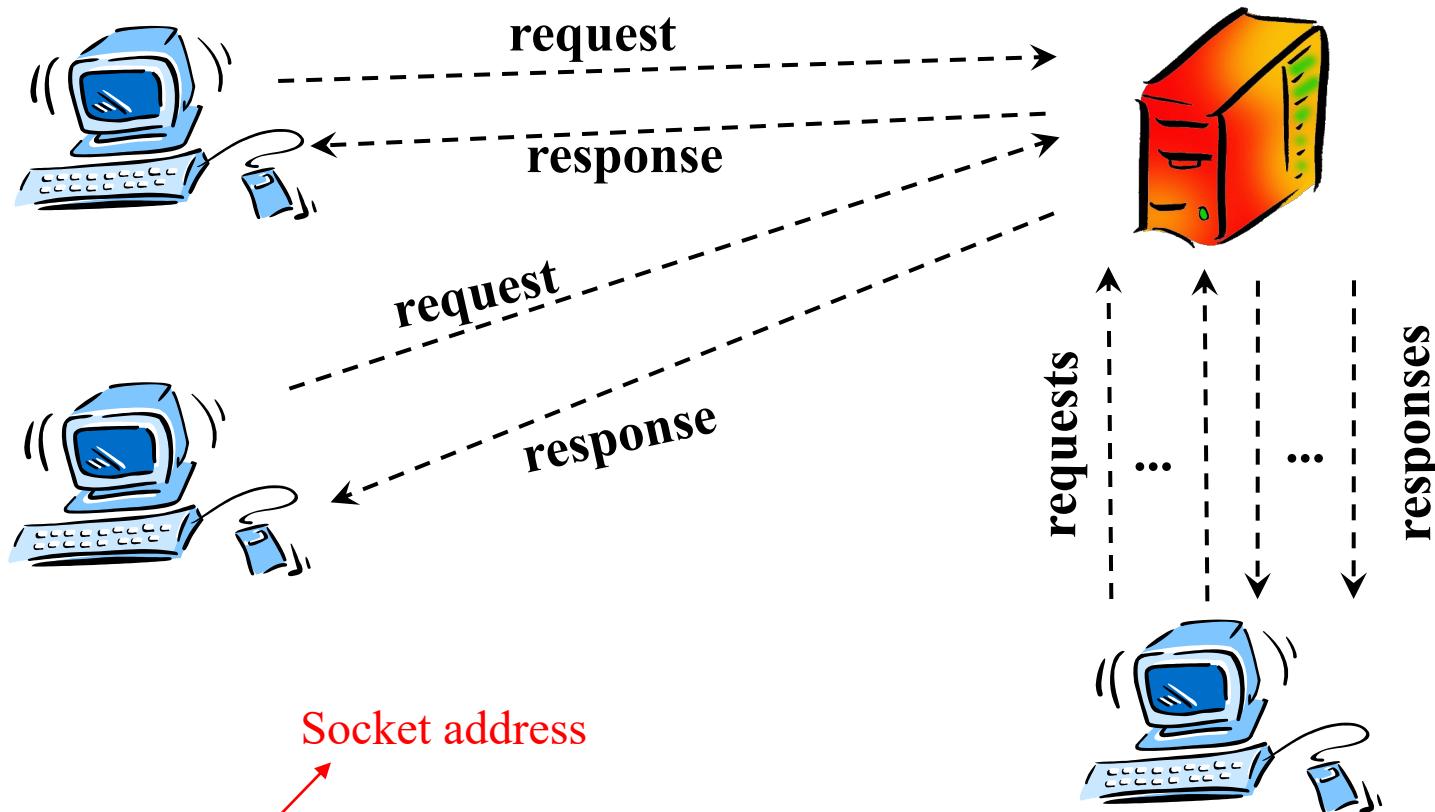
HTTP://cerbero.elet.polimi.it:8080/... (if exists!)

Port numbers

- 16 bit address (0-65535)
- well known port numbers for common servers
 - FTP 20, TELNET 23, SMTP 25, HTTP 80, POP3 110,
... (full list: RFC 1700)
- number assignment (by IANA)
 - 0 not used
 - 1-255 reserved for well known processes
 - 256-1023 reserved for other processes
 - 1024-65535 dedicated to user apps

Programmi Client e Server

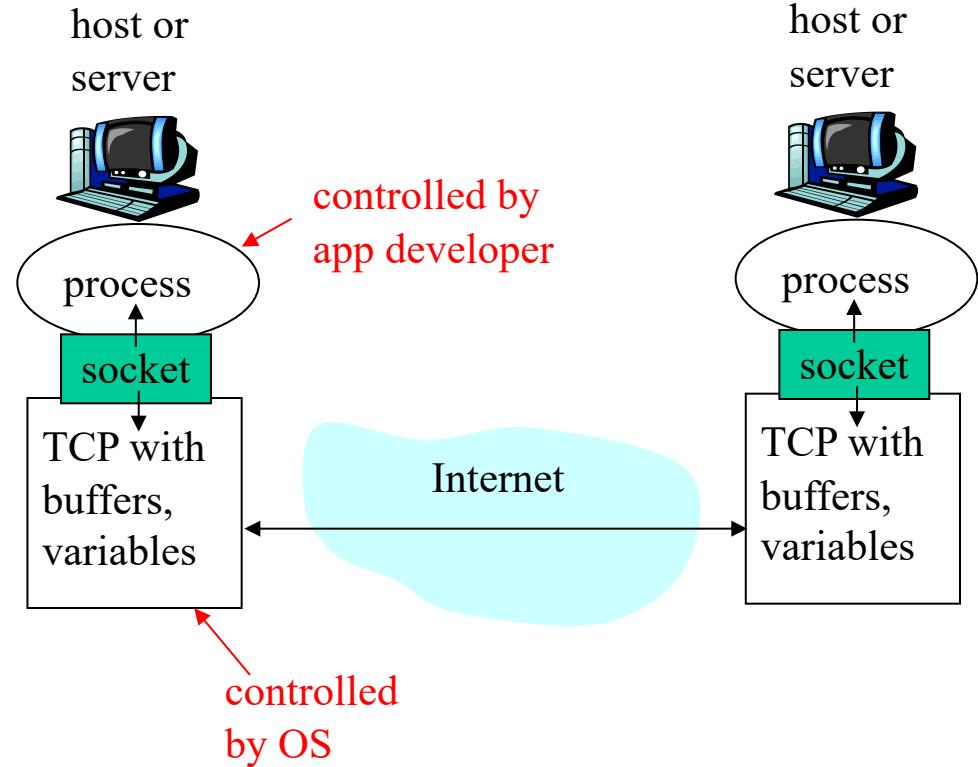
- Normalmente più client possono inviare richieste ad uno stesso server
- Un client può fare più richieste contemporanee



- Un flusso di informazioni tra due processi identificato da una quadrupla (*indirizzo IP sorgente, porta sorgente, indirizzo IP destinazione, porta destinazione*)
- Di solito il client NON USA un well known port #
 - OS assegna un numero di porta disponibile

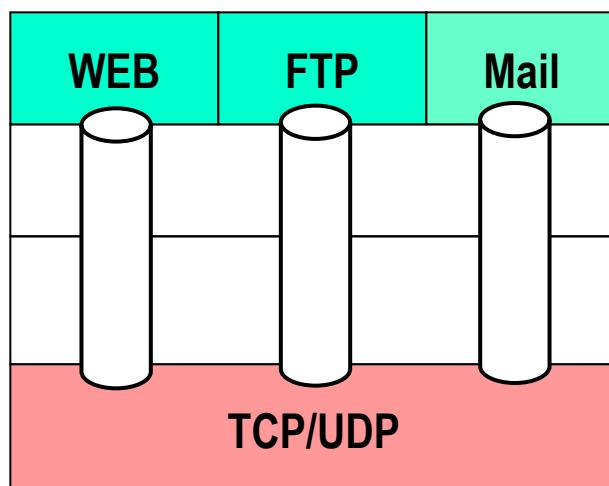
Processes communicating across network

- process sends/receives messages to/from its socket
- socket analogous to door
 - sending process shoves message out door
 - sending process assumes transport infrastructure on other side of door which brings message to socket at receiving process
- API: (1) choice of transport protocol; (2) ability to fix a few parameters (lots more on this later)



Interazione coi livelli inferiori: architettura TCP/IP

- I protocolli applicativi che si appoggiano sull' Internet Protocol si appoggiano direttamente sul protocollo di trasporto TCP/UDP
- Lo scambio di messaggi fra i processi applicativi avviene utilizzando i servizi dei livelli inferiori attraverso i SAP (Service Access Point).
- Qualsiasi protocollo di livello applicativo accede al servizio di trasporto mediante socket.
- I protocolli di trasporto estendono il servizio di trasporto in rete dell'info tra due end systems offerto da IP consentendo il trasporto dell'info tra due PROCESSI APPLICATIVI in esecuzione sui due end systems
- Come si distingue a quale protocollo di trasporto passare il segmento? Protocol number

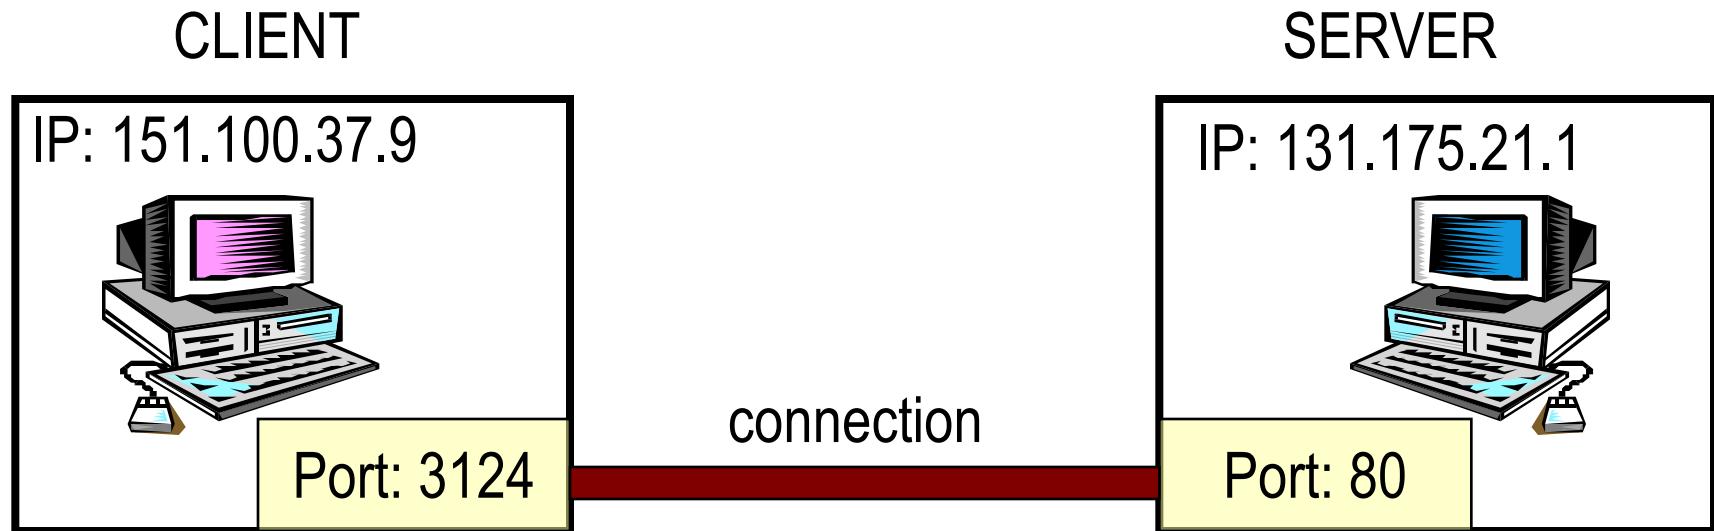


Livello controllato dall'applicazione

Livelli controllati dal sistema operativo

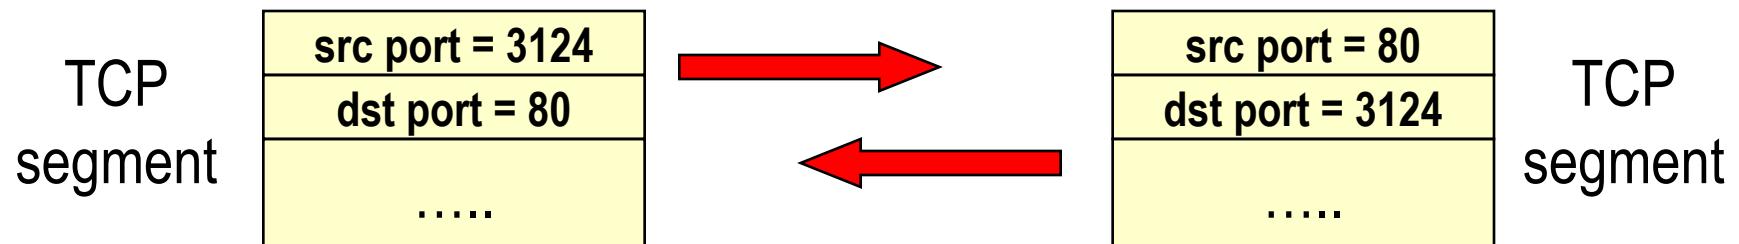
Connections

identified by socket addresses at its ends

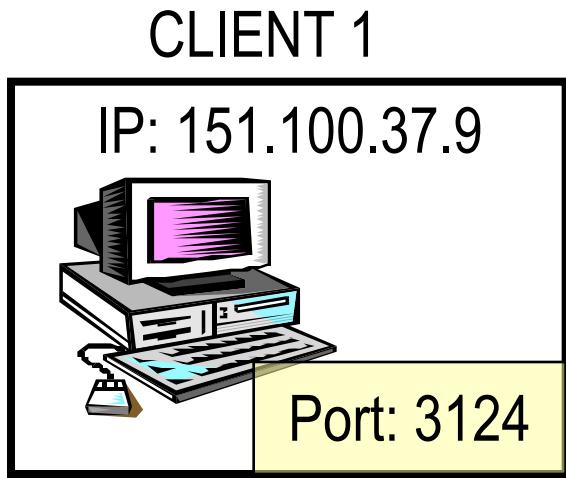


CLIENT SOCKET: [<151.100.37.9, 3124>]

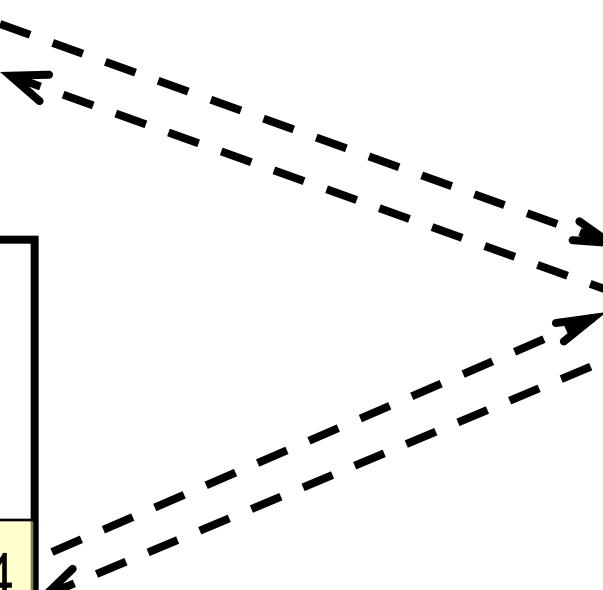
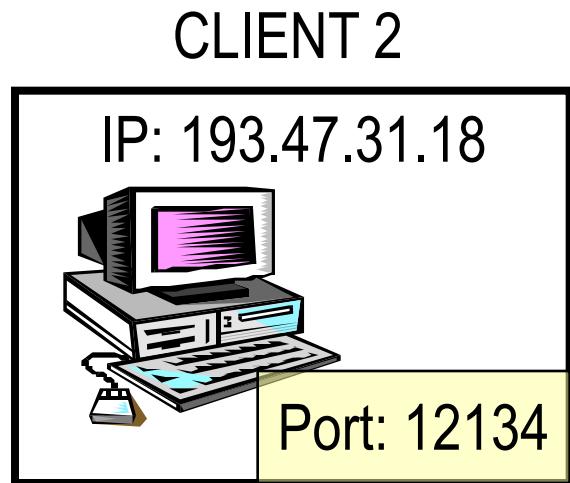
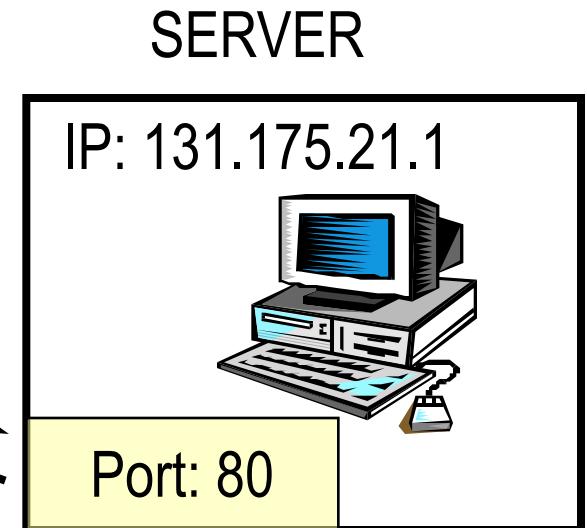
SERVER SOCKET: [<131.175.21.1, 80>]



Managing multiple connections

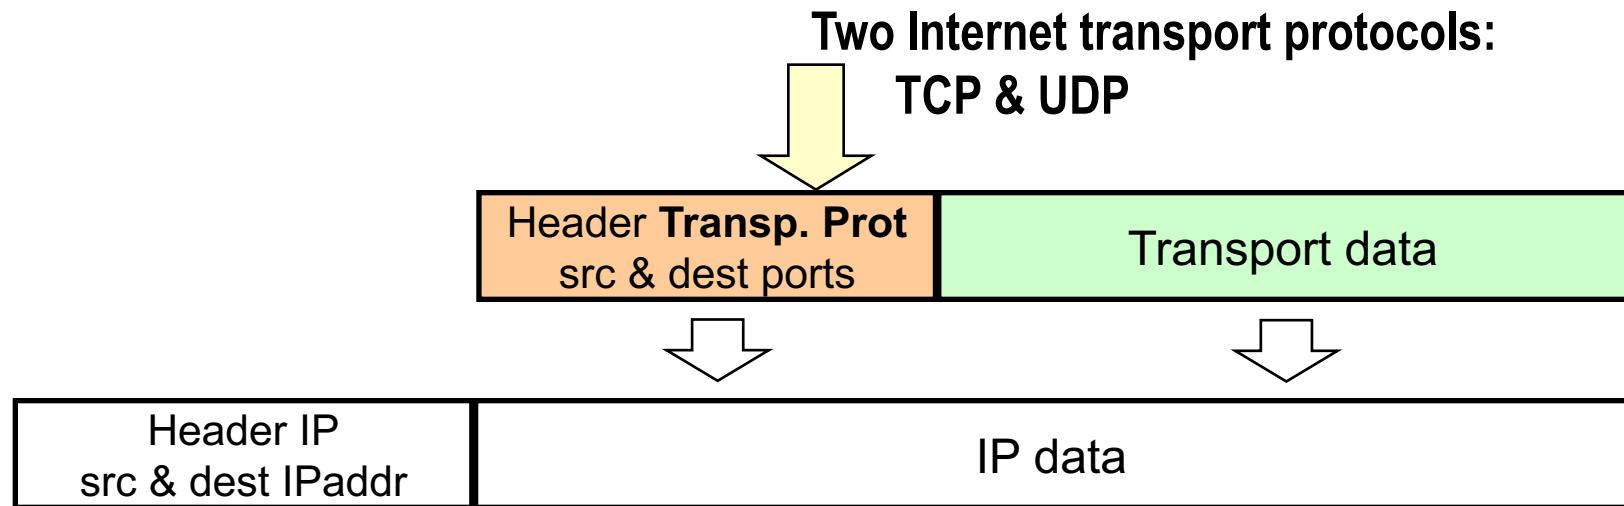


Server unique socket address can be accessed simultaneously by multiple clients



Socket addresses: where?

- ❑ Ports: in the Internet Transport protocol header (TCP or UDP)
- ❑ IP addresses: In the IP packet header



IP packets travel in the network based on IP address;

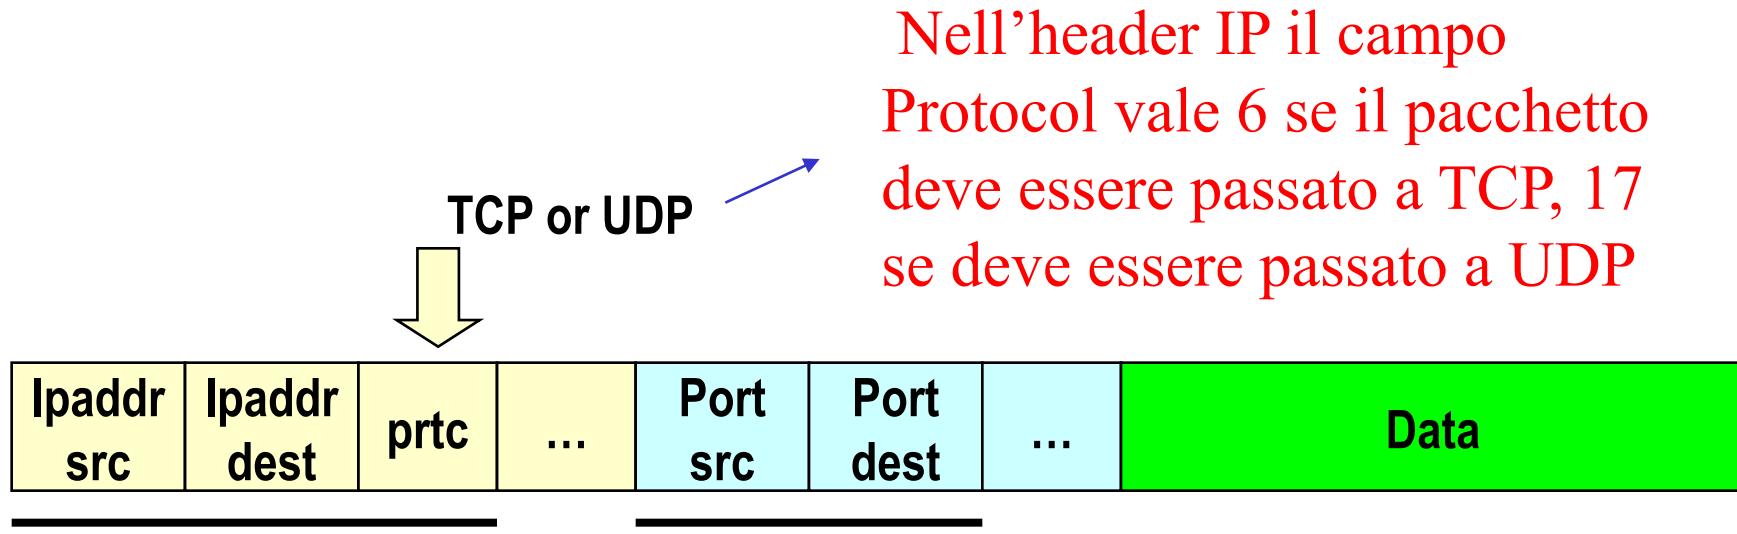
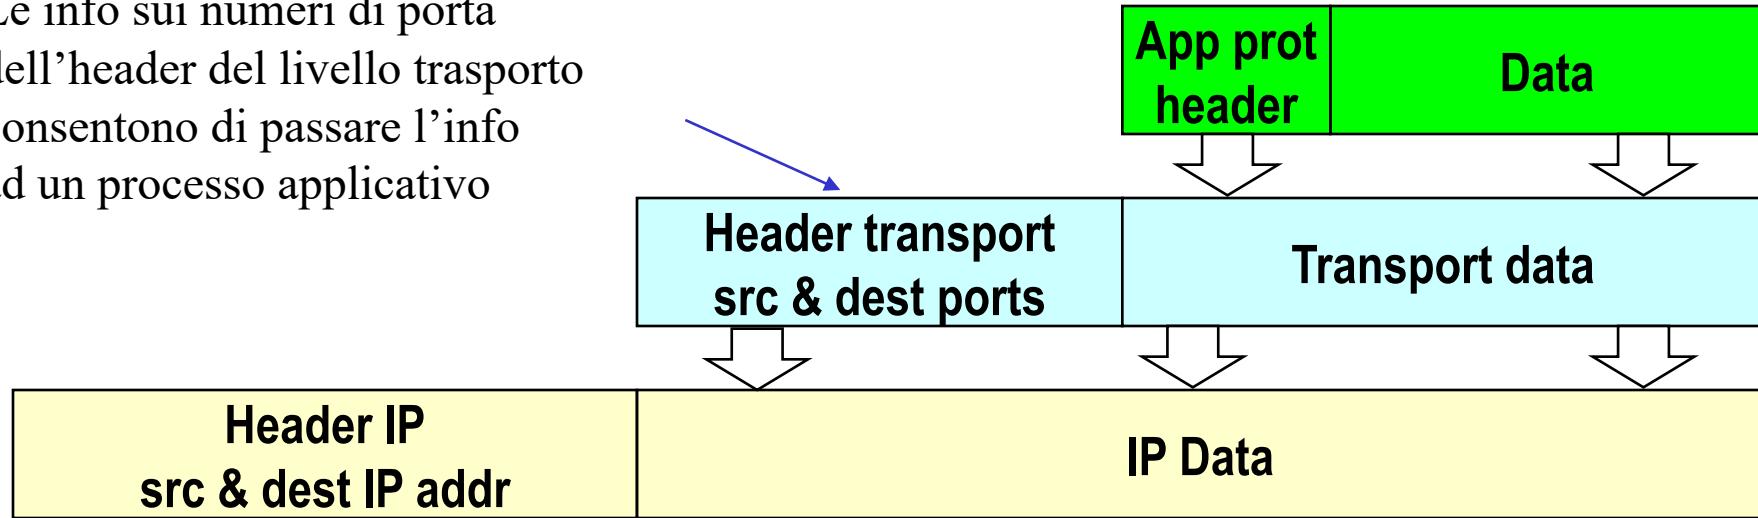
Once they reach destination, they are delivered to app based on port #

Sockets and Internet transport protocols

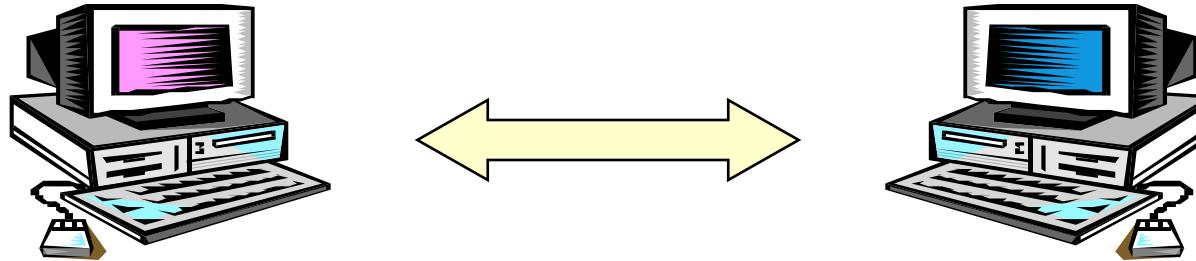
- Two transport protocols in Internet:
 - TCP (reliable, connection-oriented)
 - UDP (unreliable, connectionless)
- When opening socket, needs to specify which transport to use!

Socket addresses (refined)

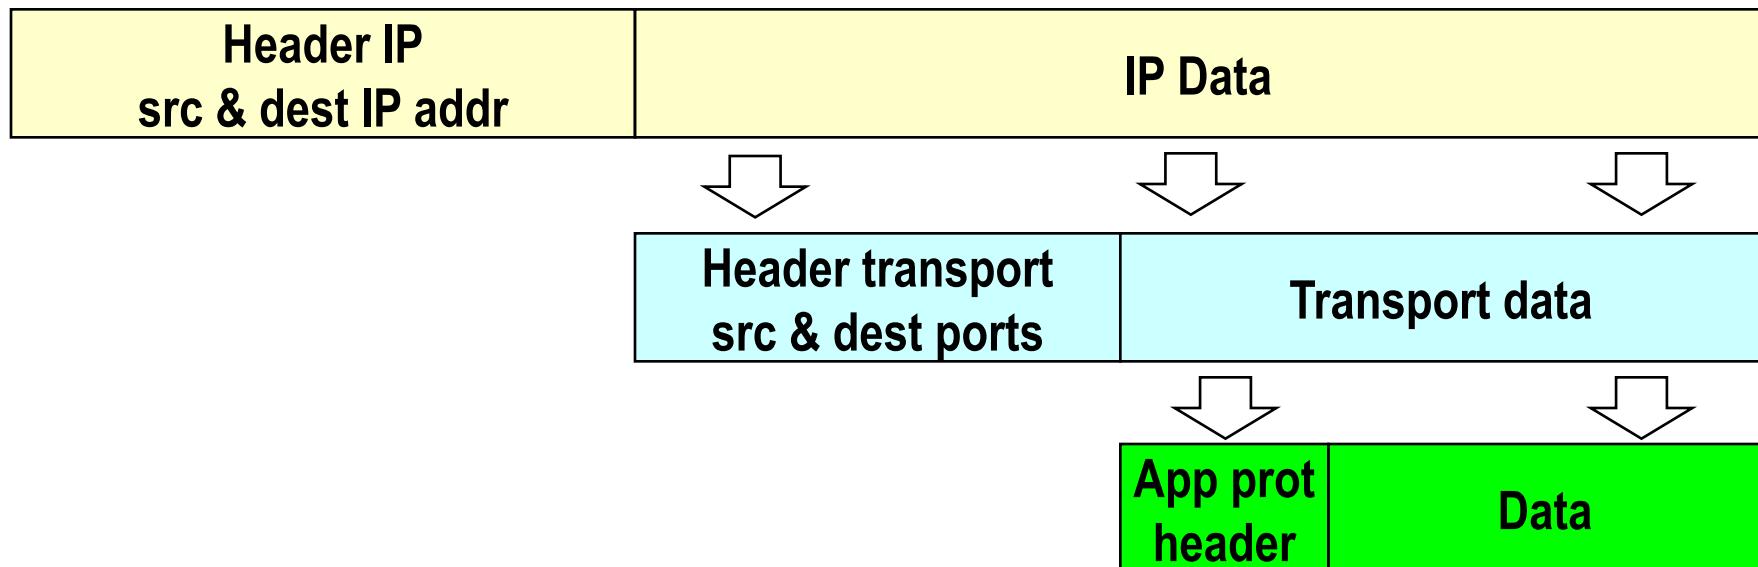
Le info sui numeri di porta
dell'header del livello trasporto
consentono di passare l'info
ad un processo applicativo



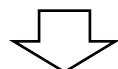
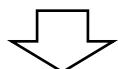
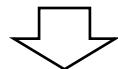
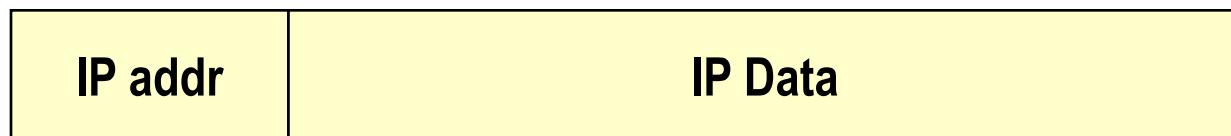
The Internet level view



Information units travelling in the network: IP packets



Demultiplexing at receiver



IP SW: checks IP packet;

Sends to transport sw

selects whether UDP or TCP

Transport demux

Transport SW: checks segment;

Sends to application sw based on Port number

Programmi Client e Server

- Un client può essere eseguito in modalità parallela o seriale
 - esempio: invia più richieste in parallelo per i file che compongono una pagina web
- Anche un server può essere eseguito in modalità parallela o seriale
- Normalmente i server che usano TCP vengono eseguiti in modalità parallela e sono dunque in grado di rispondere a più richieste contemporaneamente
- Con ognuno dei client viene aperta una connessione TCP che viene mantenuta per il tempo necessario a scambiare richieste e risposte
- La gestione delle procedure per ciascun client collegato avviene mediante la generazione di processi figli (per clonazione del processo o uso di processi multi-thread)

What transport service does an app need?

Data loss

- some apps (e.g., audio) can tolerate some loss
- other apps (e.g., file transfer, telnet) require 100% reliable data transfer

Timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

Bandwidth

- some apps (e.g., multimedia) require minimum amount of bandwidth to be “effective”
- other apps (“elastic apps”) make use of whatever bandwidth they get

Transport service requirements of common apps

Application	Data loss	Bandwidth	Time Sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video:10kbps-5Mbps	yes, 100's msec
stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps up	yes, 100's msec
instant messaging	no loss	elastic	yes and no

Internet transport protocols services

TCP service:

- ❑ *connection-oriented*: setup required between client and server processes
- ❑ *reliable transport* between sending and receiving process
- ❑ *flow control*: sender won't overwhelm receiver
- ❑ *congestion control*: throttle sender when network overloaded
- ❑ *does not providing*: timing, minimum bandwidth guarantees

UDP service:

- ❑ unreliable data transfer between sending and receiving process
- ❑ does not provide: connection setup, reliability, flow control, congestion control, timing, or bandwidth guarantee

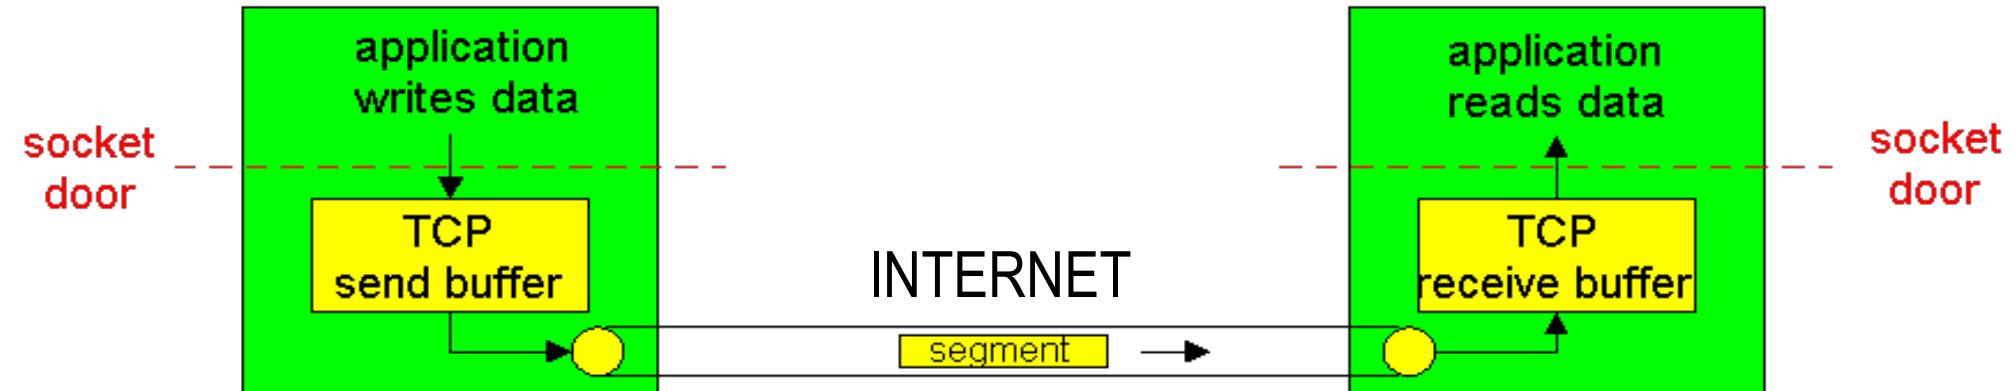
Q: why bother? Why is there a UDP? Which service does it add to IP?

Multiplexing/Demultiplexing
Error Control

Internet apps: application, transport protocols

Application	Application layer protocol	Underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	proprietary (e.g. RealNetworks)	TCP or UDP
Internet telephony	proprietary (e.g., Dialpad)	typically UDP

Why it is trivial (!) to write networking apps?



- Application software duties:
 - open socket (e.g. C, C++, JAVA function call, OS call, external library primitive)
 - Injects message in its own socket
 - being confident message is received on the other side
- TCP software: in charge of managing segments!
 - reliable message transport when TCP used
 - Segmentation performed by TCP transmitter
 - Receive buffer necessary to ensure proper packet's order & reassembly

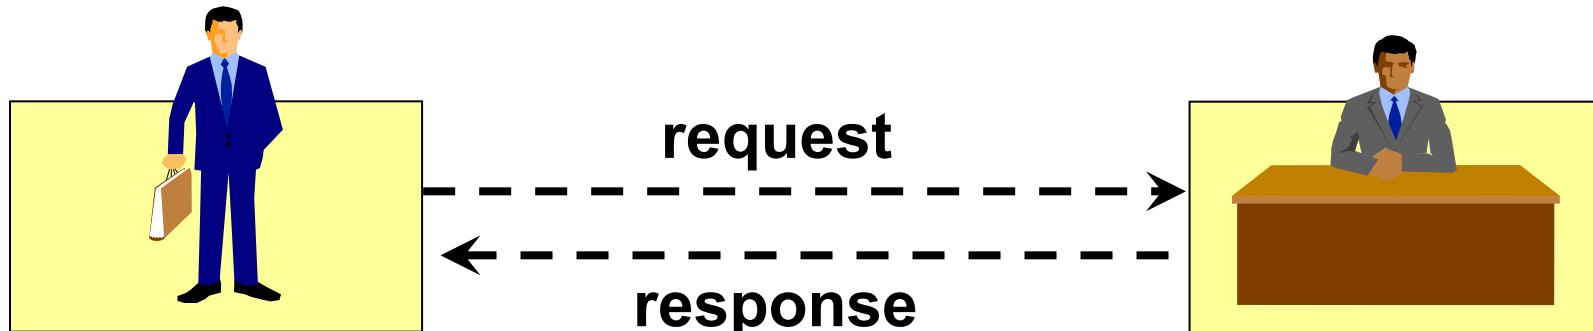
1. Web Application: HTTP Protocol

□ HTTP: the protocol of the WWW

- version 1.0, RFC 1945, may 1996
- version 1.1, RFC 2068 (jan97), RFC 2616 (jun99)

(but also FTP: file transfer protocol, TELNET: opens a telnet window, FILE: access local file etc.)

□ HTTP: Request-response protocol. A request message is sent from the client to a recipient server. The recipient server sends a response message back.



□ Different requests for the same URL can result in different responses due to the time of request, changes in the resource, the request header fields.

The three components of an URL

1. Protocol (also called "scheme")

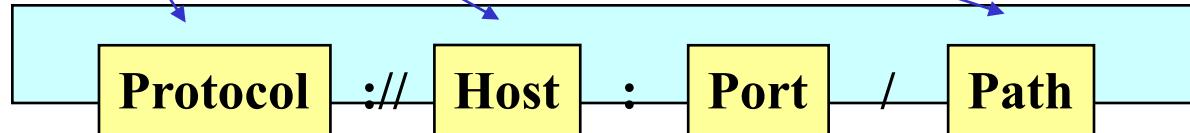
- how can a page be accessed? (application protocol used)
 - **http://www.dsi.uniroma1.it/people/petrioli/index.html**

2. Host name

- Where is the page located? (symbolic or numeric location)
 - **http://WWW.dsi.uniroma1.it/people/petrioli/index.html**

3. File (resource) name

- What is the page called? (with full path)
 - **http://www.dsi.uniroma1.it/people/petrioli/index.html**



2. Location - host name

Specifies where is the page located:

➤ **on which host**

- Humans understand names;
- Machines prefer numbers!
- Domain Name System (DNS) protocol:
 - translates names in numbers

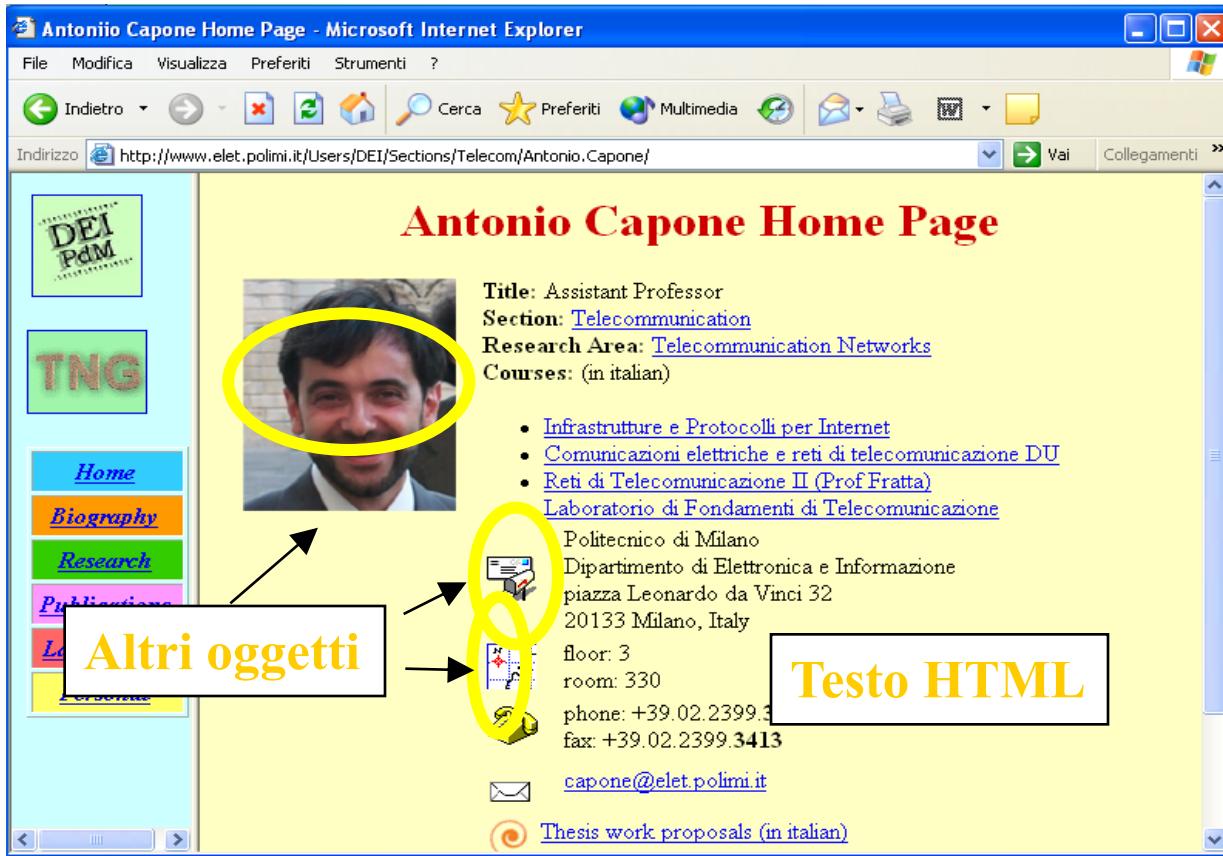
www.elet.polimi.it

DNS

131.175.21.1

Trasporto dei messaggi

- Supponiamo che un client richieda una pagina HTML di un server al cui interno sono contenuti i riferimenti ad altri oggetti (ad esempio 10 figure che compongono la pagina e che occorre visualizzare insieme al testo HTML).



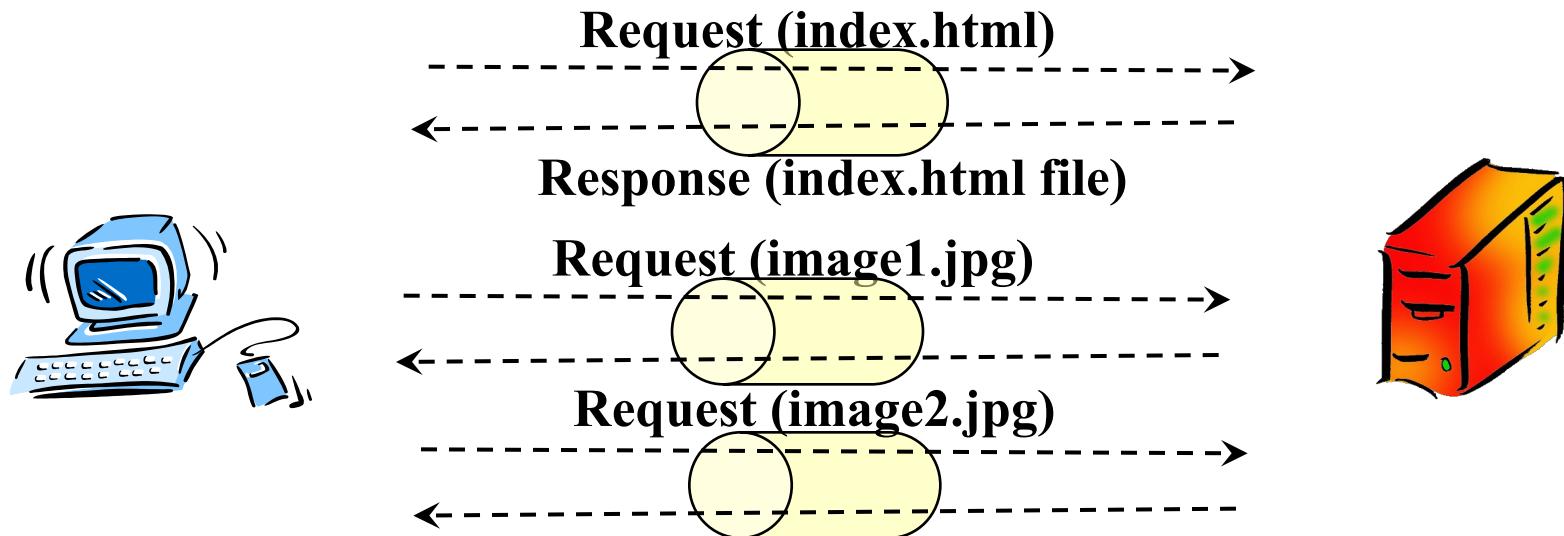
■ nel trasferimento dell'insieme di oggetti sono possibili 2 modalità

■ Non-persistent connection (default mode di HTTP 1.0)

■ Persistent connection (default mode di HTTP 1.1)

Non persistent

- Viene aperta una connessione per una sola request-response: inviato l'oggetto, il server chiude la connessione
- La procedura viene ripetuta per tutti i file collegati al documento HTML base



Nonpersistent HTTP (HyperText Transfer Protocol)

Suppose user enters URL

www.someSchool.edu/someDepartment/home.index

(contains text,
references to 10
jpeg images)

1a. HTTP client initiates TCP connection to HTTP server (process) at www.someSchool.edu on port 80

1b. HTTP server at host www.someSchool.edu waiting for TCP connection at port 80. “accepts” connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object someDepartment/home.index

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time ↓

Nonpersistent HTTP (cont.)

time
↓

4. HTTP server closes TCP connection.

5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects
6. Steps 1-5 repeated for each of 10 jpeg objects

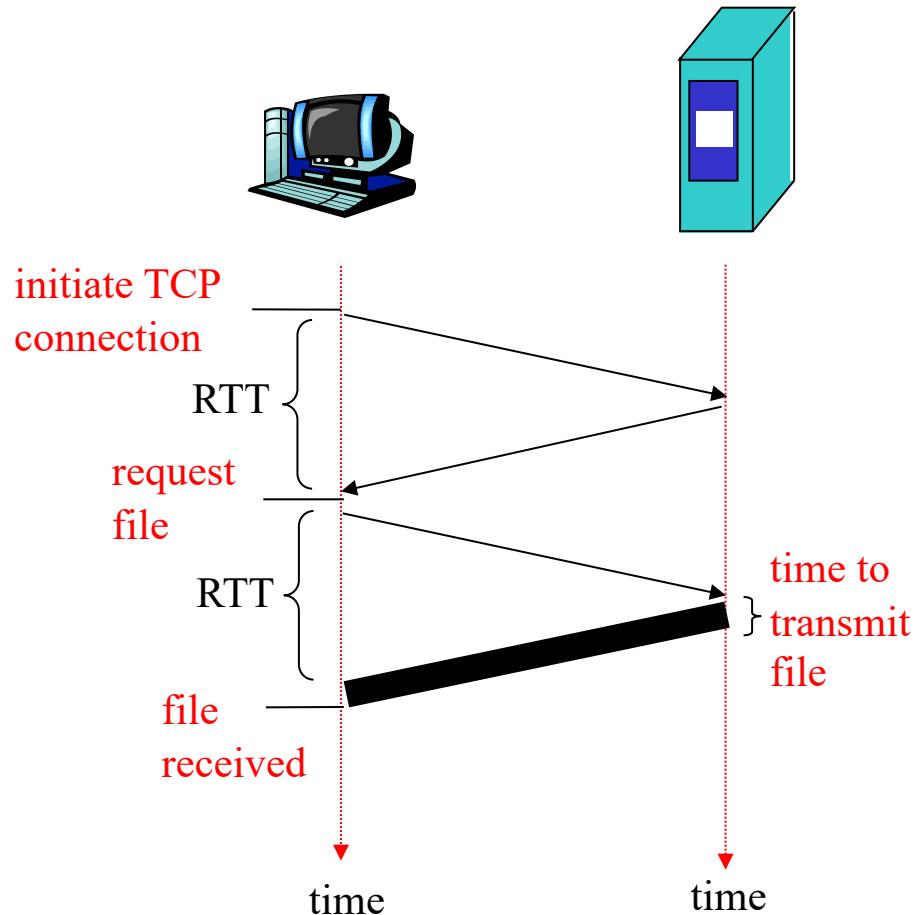
Response time modeling

Definition of RRT: time to send a small packet to travel from client to server and back.

Response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time

total = 2RTT+transmit time



Problema: OGNI VOLTA CHE SI APRE UNA CONNESSIONE SI PAGA 2RTT PRIMA DI SCAMBIARE DATI

Non persistent

- Viene aperta una connessione per una sola request-response: inviato l'oggetto, il server chiude la connessione
- La procedura viene ripetuta per tutti i file collegati al documento HTML base
- Le connessioni TCP per più oggetti possono essere aperte in parallelo per minimizzare il ritardo
- Il numero max di connessioni è di solito configurabile nel browser



Un esempio...

- Esempio: di richiesta oggetto

```
GET /ntw/index.html HTTP/1.0
Date: Wed, 22 Mar 2000 09:09:01 GMT
Pragma: No-cache
From: gorby@moskvax.com
User-agent: Mozilla/4.3
```

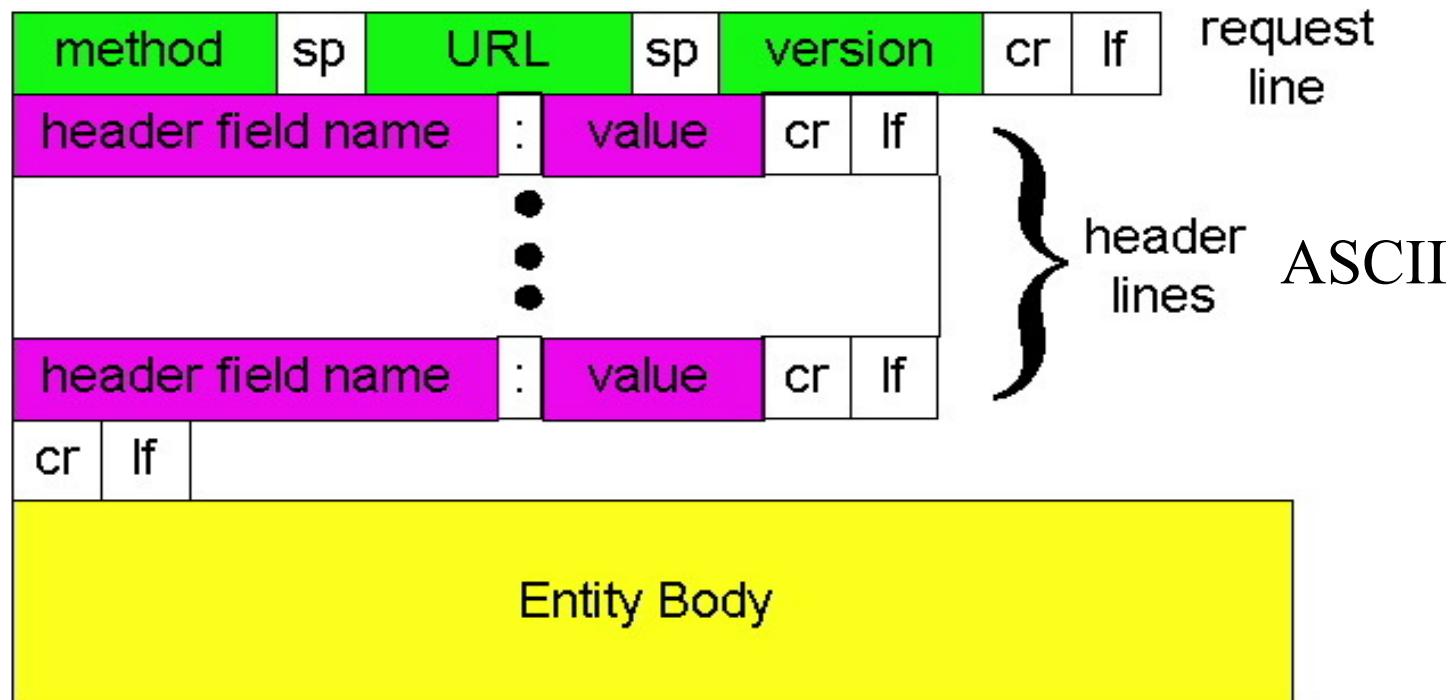
HTTP è testuale (ASCII)

- Esempio: risposta

```
HTTP/1.0 200 OK
Date: Wed, 22 Mar 2000 09:10:01 GMT
Server: Apache/1.3.0 (Unix)
Content-Length: 6821
Last-Modified: Mon, 22 Jun 1998 09:23:24 GMT
Content-Type: text/html
data data data data data ...
```

HTTP Messages

Request



HTTP Message

Methods:

GET	E' usato quando il client vuole scaricare un documento dal server. Il documento richiesto è specificato nell'URL. Il server normalmente risponde con il documento richiesto nel corpo del messaggio di risposta.
HEAD	E' usato quando il client non vuole scaricare il documento ma solo alcune informazioni sul documento (come ad esempio la data dell'ultima modifica). Nella risposta il server non inserisce il documento ma solo degli header informativi.
POST	E' usato per fornire degli input al server da utilizzare per un particolare oggetto (di solito un applicativo) identificato nell'URL.
PUT	E' utilizzato per memorizzare un documento nel server. Il documento viene fornito nel corpo del messaggio e la posizione di memorizzazione nell'URL.

■ Altri methods:

■ DELETE, LINK, UNLINK, ...

Per inviare info al server

E' possibile anche usare GET

GET /search.cgi?string=greek-architects HTTP/1.0

Un esempio...

- Esempio: di richiesta oggetto

```
GET /ntw/index.html HTTP/1.0
Date: Wed, 22 Mar 2000 09:09:01 GMT
Pragma: No-cache
From: gorby@moskvax.com
User-agent: Mozilla/4.3
```

HTTP è testuale (ASCII)

- Esempio: risposta

```
HTTP/1.0 200 OK
Date: Wed, 22 Mar 2000 09:10:01 GMT
Server: Apache/1.3.0 (Unix)
Content-Length: 6821
Last-Modified: Mon, 22 Jun 1998 09:23:24 GMT
Content-Type: text/html
data data data data data ...
```

Header

Header name : **Header value**

- Gli header servono per scambiare informazione di servizio aggiuntiva
- E' possibile inserire più linee di header per messaggio
- Esempi

Cache-control	Informazione sulla cache
Accept	Formati accettati
Accept-language	Linguaggio accettato
Authorization	Mostra i permessi del client
If-modified-since	Invia il doc. solo se modificato
User-agent	Tipo di user agent

Header Types

- General: used in request and response messages
- Request header (e.g. to express preferences on the nature of the response, include additional info with the request, to specify a constraint on the server in handling the request)
- Response header: to provide additional info about the response or to request additional info from the user
- Entity header (both in request and response messages) to provide info on the entity such as the last time it was modified

Header Types-Request

- General: used in request and response messages

Date	Indicates the date and time of the message origination, e.g. Date: Tue 16 May 2000 11:29:32 GMT
Pragma	Permits to send directives to the recipient, requesting it to behave in a particular way while handling a request or response. Pragma: no-cache informs proxies on the path not to return a cached copy

- Request header (e.g. to express preferences on the nature of the response, include additional info with the request, to specify a constraint on the server in handling the request)

Authorization	Includes credentials required to access a resource (e.g. Authorization: Basic YXZpYXRpS29IDizDizM1NA== where Basic is the authentication scheme and YXZpYXRpS29IDizDizM1NA== an encoding of user id and password)
From	The user may include e-mail address for identification (From: gorby@moskvax.com)
If-Modified-Since	Indicates not to return a copy of the resource if it has not been modified after teh specified date. Used for caching at a proxy or browser (saves the time needed for downloading of the resource). GET /foo.html HTTP/1.0 If-Modified-Since: Sun, 21 May 200 07:00:25 GMT
Referer	Includes the URL from which the requested URL was obtained. Can be used for locating obsolete links. Referer: http://www.cnn.com
User-Agent	User browser used, client machine OS, etc User-Agent: Mozilla/4.04 [en]C-WorldNet (win95;l)

Header Types

- Entity header (both in request and response messages) to provide info on the entity such as the last time it was modified

Allow	Indicates the list of valid methods that can be applied to a resource PUT /foo.html HTTP/1.0 Allow: HEAD, GET, PUT
Content-Type	Media type of the entity body (e.g. image/gif, text/html etc.)
Content-Encoding	Indicates how the resource representation can be decoded into the format indicated in the Content-Type field (e.g. if the file has been compressed with gzip) Content-Encoding:gzip
Content-Length	Length of the entity body (bytes). Allow to check whether all the body was received
Expires	The entity should be considered stale after the time specified in this header (after this time the content could still be stored in cache but should not be returned to users without first validating with the origin server)
Last-Modified	Specify the time at which the resource was modified last

Un esempio...

- Esempio: di richiesta oggetto

```
GET /ntw/index.html HTTP/1.0
Date: Wed, 22 Mar 2000 09:09:01 GMT
Pragma: No-cache
From: gorby@moskvax.com
User-agent: Mozilla/4.3
```

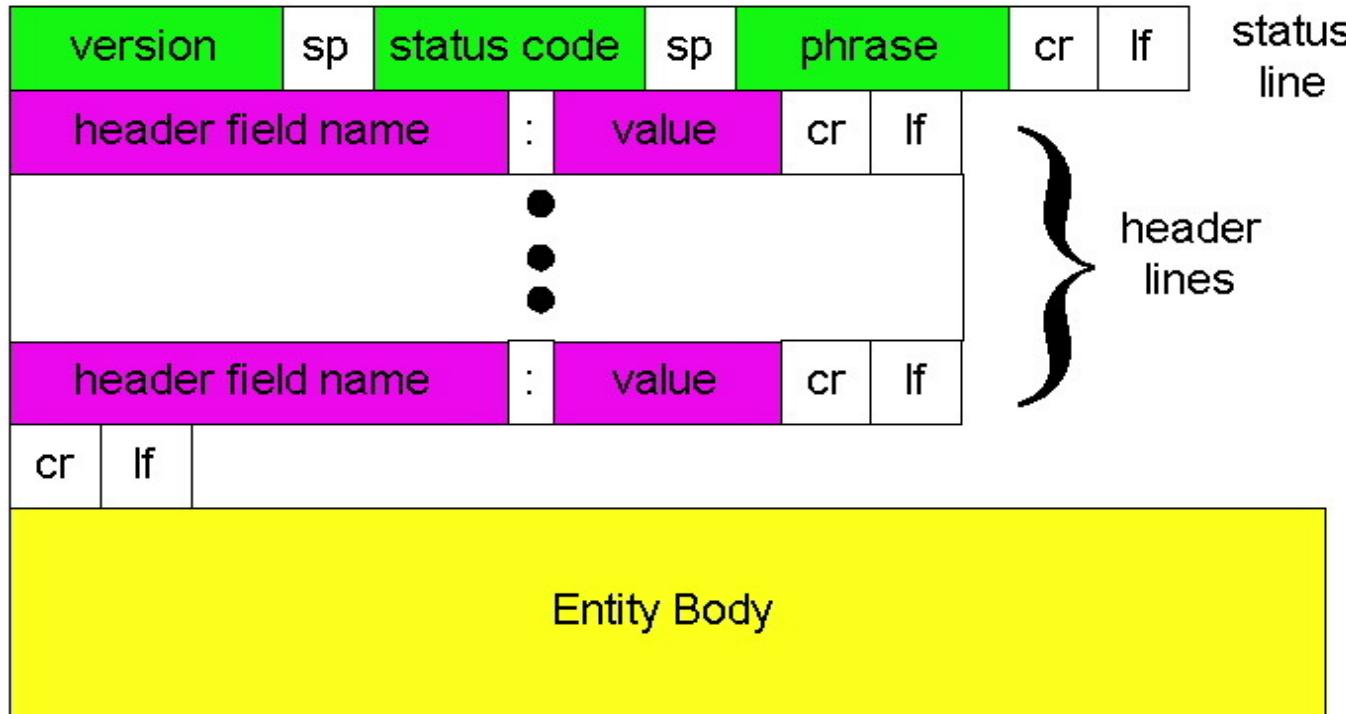
HTTP è testuale (ASCII)

- Esempio: risposta

```
HTTP/1.0 200 OK
Date: Wed, 22 Mar 2000 09:10:01 GMT
Server: Apache/1.3.0 (Unix)
Content-Length: 6821
Last-Modified: Mon, 22 Jun 1998 09:23:24 GMT
Content-Type: text/html
data data data data data ...
```

Message

Response



Messaggi

■ 1xx Informational

Status codes:

■ 2xx Success

200 OK:	La richiesta ha avuto successo; l'informazione è inclusa
201 Created:	La risorsa e' stata creata con successo in seguito ad UN POST (se non puo' essere creata subito 202 Accepted)
202 Accepted:	The request has been received but has not yet been handled in full (e.g., since a program must run which requires a long time → the user agent can continue with its task without waiting for the action to complete at the origin server)
204 No content:	Successfully received, no change required in what the user actually sees

Messaggi

■ 3xx Redirection

301 Moved Permanently:	L'oggetto è stato spostato nell'URL indicato
302 Moved Temporarily:	L'oggetto è stato spostato nell'URL indicato
304 Not Modified:	L'oggetto non modificato dal tempo incluso nella richiesta

■ 4xx Client error

400 Bad Request:	errore generico (sintassi errata/irriconoscibile)
401 Unauthorized:	Autenticazione fallita (e.g., Accesso senza necessari account e password)
403 Forbidden:	La richiesta e' stata ricevuta e compresa ma il server ha stabilito di non servirla. Le ragioni possono essere indicate nell'entity body.
404 Not Found:	l'oggetto non esiste sul server

■ 5xx Server error

500 Internal server error	Generico. Errore o guasto nel server
501 Not implemented	Funzione non implementata
503 Service unavailable	Servizio non disponibile

Header Types-Response

- Response header: to provide additional info about the response or to request additional info from the user

Location	Used to redirect the request to where the resource can be found Location: http://www.foo.com/level1/twosdown/Location.html
Server	Origin server SW version number and configuration related info Server: Apache/1.2.6 Red Hat
WWW-Authenticate	Request to retransmit the request with appropriate credentials according to a given scheme

Un esempio

□ Esempio: di richiesta oggetto

```
PUT /motd HTTP/1.0
Date: Wed, 22 Mar 2000 08:10:07 GMT
From: gorby@moskvax.com
User-agent: Mozilla/4.0
Content-length:23
Allow:GET,HEAD,PUT
data data data data data ...
```

HTTP è testuale (ASCII)

□ Esempio: risposta

```
HTTP/1.1 200 OK
Date: Wed, 22 Mar 2000 09:10:07 GMT
Server: Apache/1.3.0 (Unix)
...
```

Cache locale: get condizionato

- E' possibile evitare di scaricare oggetti memorizzati nella memoria locale se non sono stati modificati

Client:

GET /fruit/kiwi.gif HTTP/1.0

User-agent: Mozilla/4.0

Accept: text/html, image/gif, image/jpeg

If-modified-since: Mon, 22 Jun 1998 09:23:24

Server:

HTTP/1.0 304 Not Modified

Date: Wed, 19 Aug 1998 15:39:29

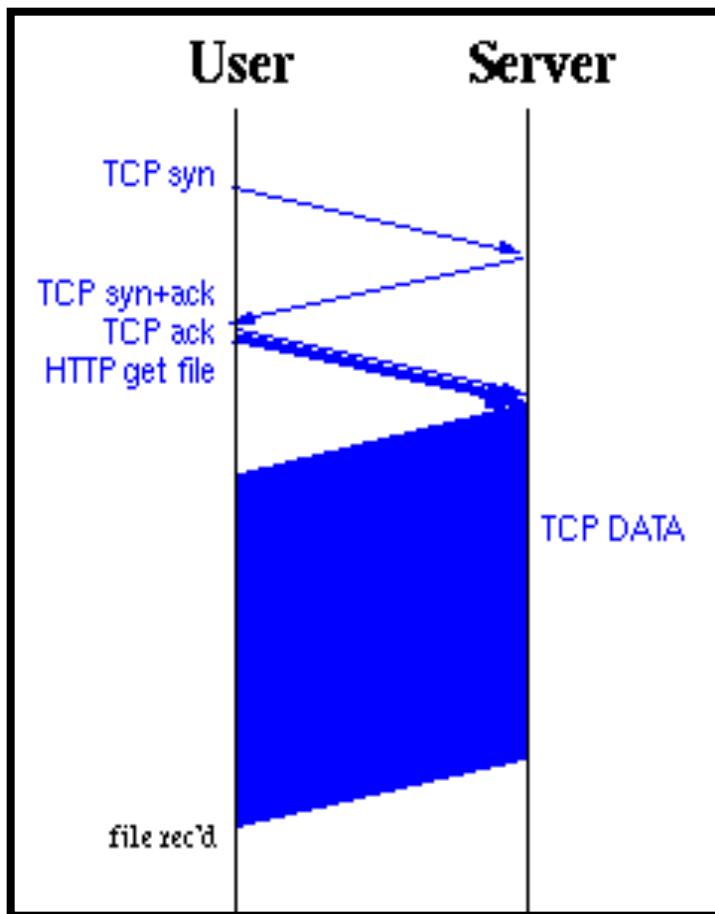
Server: Apache/1.3.0 (Unix)

(empty entity body)

- E' possibile anche usare il metodo HEAD

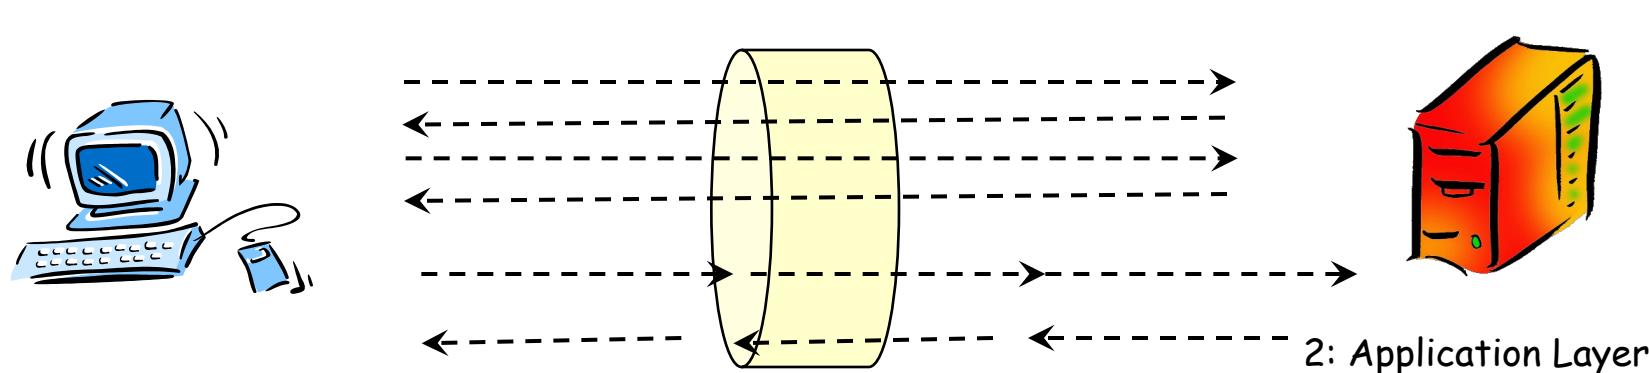
Performance drawbacks

- Mandatory roundtrips
 - TCP three-way handshake
 - get request, data return
 - new connections for each inlined image (parallelize)



Persistent connection (HTTP v1.1)

- Nel caso *persistent* il server non chiude la connessione dopo l'invio dell'oggetto
- La connessione rimane aperta e può essere usata per trasferire altri oggetti della stessa pagina web o anche più pagine
- I server chiudono di solito la connessione sulla base di un *time-out*
- *without pipelining*: il client invia una nuova richiesta solo dopo aver ricevuto la risposta per la precedente
 - *with pipelining*: più richieste vengono inviate consecutivamente dal client



Persistent HTTP

Nonpersistent HTTP issues:

- requires 2 RTTs per object
- OS must work and allocate host resources for each TCP connection
- but browsers often open parallel TCP connections to fetch referenced objects

Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server are sent over connection

Persistent without pipelining:

- client issues new request only when previous response has been received
- one RTT for each referenced object

Persistent with pipelining:

- default in HTTP/1.1
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

HTTP 2.0 what's new

Performance optimization is at the core of HTTP evolution

- Streams (bidirectional flow of information) multiplexing
- Flow control at the stream level
- Prioritization
- Header compression
- Security, encryption

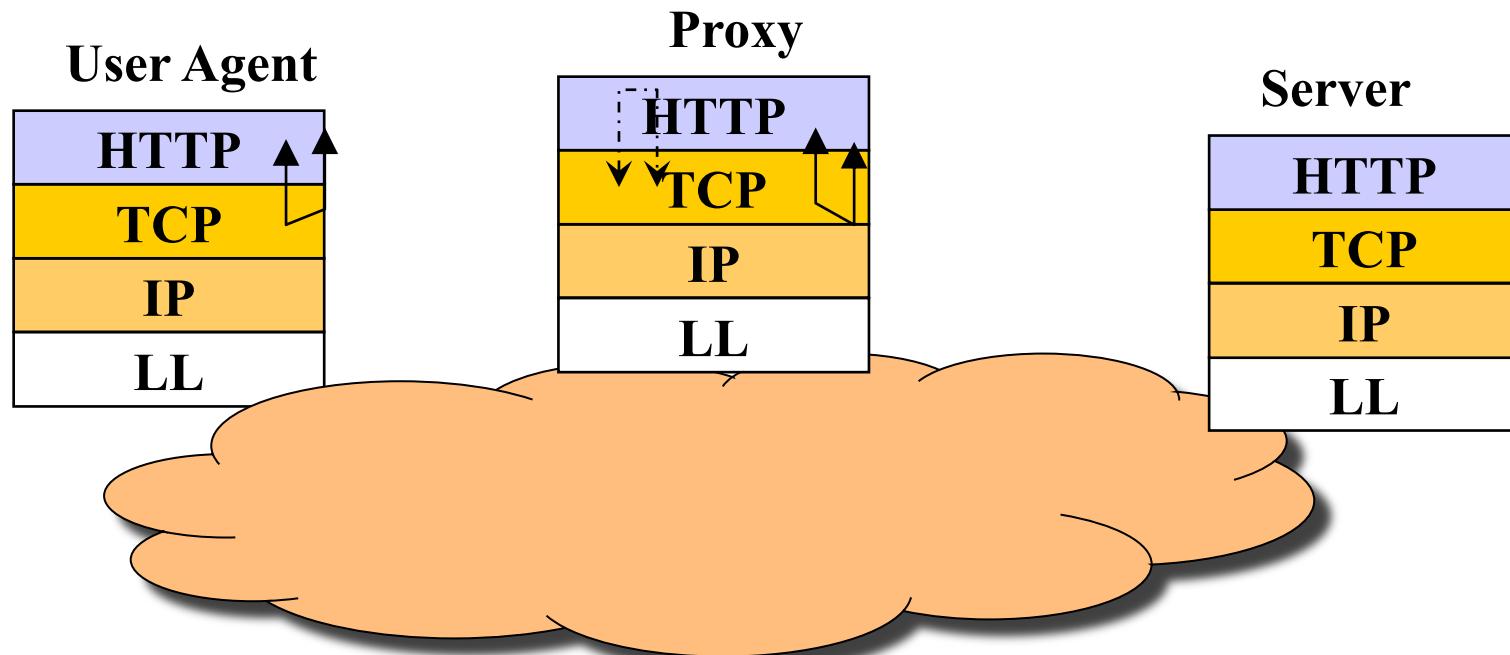
Want to know more?

<https://datatracker.ietf.org/doc/html/rfc9113#name-field-section-compression-a>

Proxy

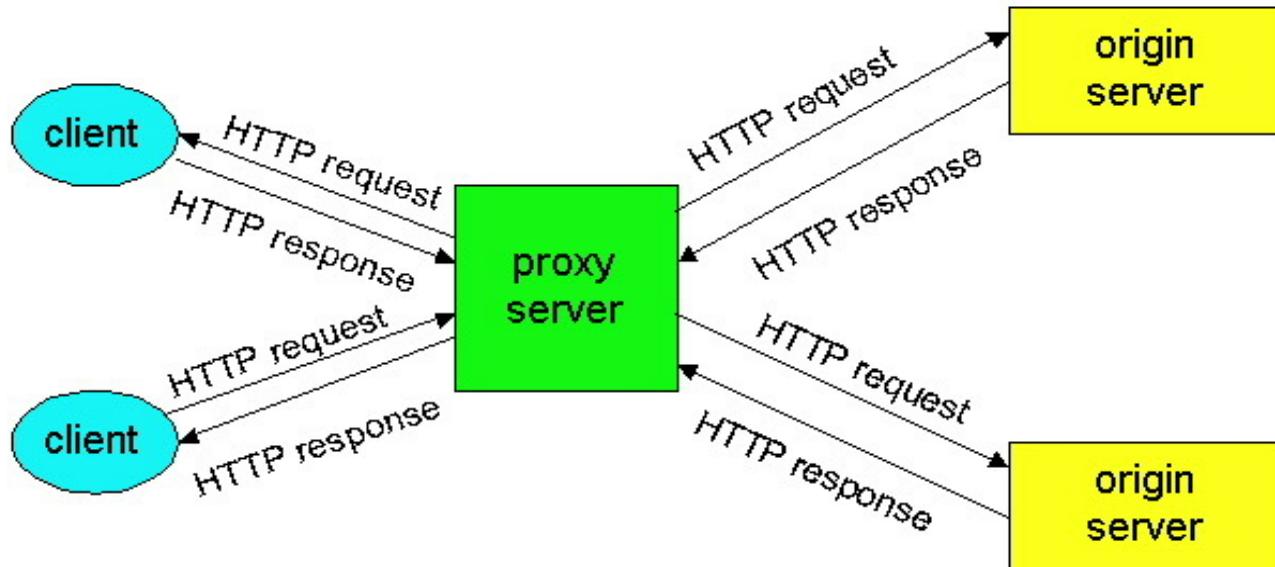
An intermediary program which acts as both a server and a client for the purpose of making requests on behalf of other clients.

- I proxy sono degli instradatori di messaggi di livello applicativo
- Devono essere sia client che server
- Il server vede arrivare tutte le richieste dal proxy (mascheramento degli utenti del proxy)



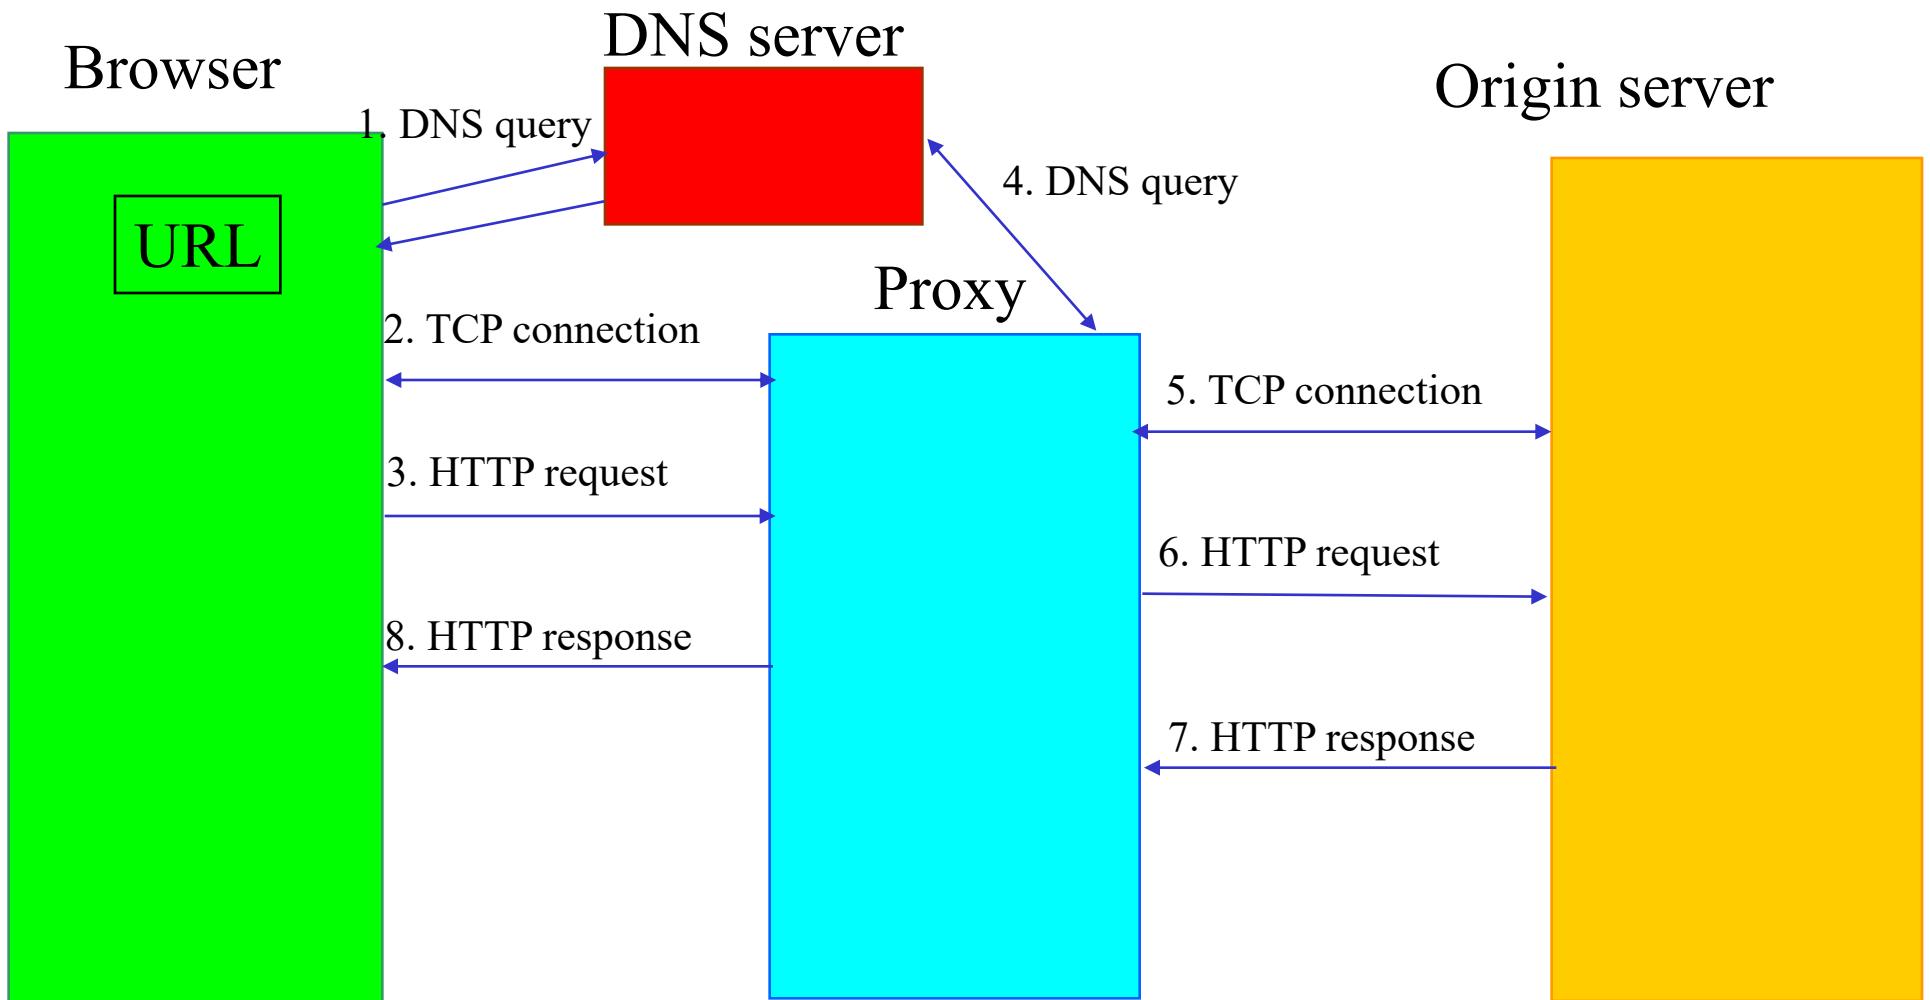
Cache di rete: uso dei proxy

- Compito principale dei proxy è fornire una grande memoria di cache
- Se un documento è contenuto nella cache viene scaricato più velocemente sul client
- Razionale: i server forniscono le stesse info o info fortemente ridondanti a comunità locali di utenti



- Consistency: cached pages might not be the most recent...
- Ad clickthrough counts: how does Yahoo know how many times you accessed their pages, or **more importantly**, their ads?

Download attraverso un proxy...



Tipi di proxy

- Regular proxies
 - forwards requests and responses
- Caching proxy
 - has a cache of responses received in the past
 - when the proxy receives a request that can be satisfied by the cache the request message is not forwarded and the response is returned directly by the proxy
- Forwarding behavior
 - Transparent proxy
 - Only superficially changes the request (e.g., adding its ID)
 - Non transparent proxy
 - Can modify request/response
 - Anonymization
 - Media type transformation (e.g. changing image format to reduce size)
 - Language transformation ...
 - Entrambi i transparent/non-transparent proxy possono avere associata una cache

Uso dei Proxy(1)

- Caching: '*storage of a response obtained earlier for later use, when clients request the same resource*'. Se la pagina richiesta e' nella cache il proxy puo' inviarla immediatamente al client purche' sia fresh (ovvero non sia stata modificata rispetto alla copia in cache)
 - latency minore MA cache consistency importante. Strong consistency/weak consistency
- Agisce come front-end per un gruppo di utenti che condividono l'accesso al web (caching ma anche suddivisione delle risorse—pro: se piu' utenti richiedono la stessa risorsa ad un server un'unica connessione deve essere stabilita; -- con: possibili ritardi dovuti al fatto che piu' utenti condividono le stesse risorse)

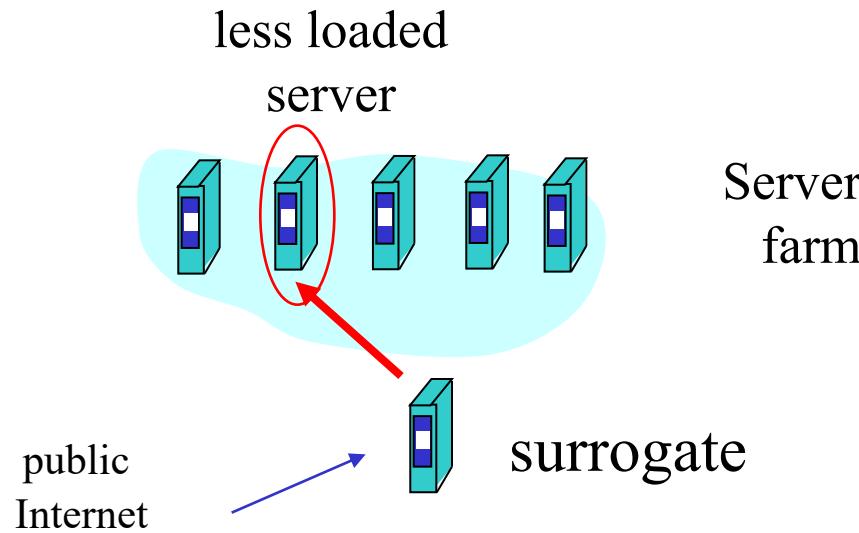
Uso dei Proxy(2)

- Mascheramento degli utenti del proxy: Il server vede arrivare tutte le richieste dal proxy e quindi non e' in grado di riconoscere l'utente. Il proxy conosce le informazioni di accesso relative agli utenti ma e' considerata un'entita' trusted.
 - Alcuni proxy non permettono l'anonymizzazione e indicano nell'header informazioni sul client
 - Mascheramento non completo: e.g., user agent e cookies
- Customizzazione delle risposte
- Filtraggio delle richieste e delle risposte
 - per limitare l'accesso a certi siti
 - filtrare richieste a siti non adatti ai minori
 - verificare la presenza di virus prima di inviare al client le risposte
 - togliere alcune informazioni dall'header prima di inviare la richiesta (es. e-mail dell'utente)

Other types of proxies

□ Reverse proxy (surrogates)

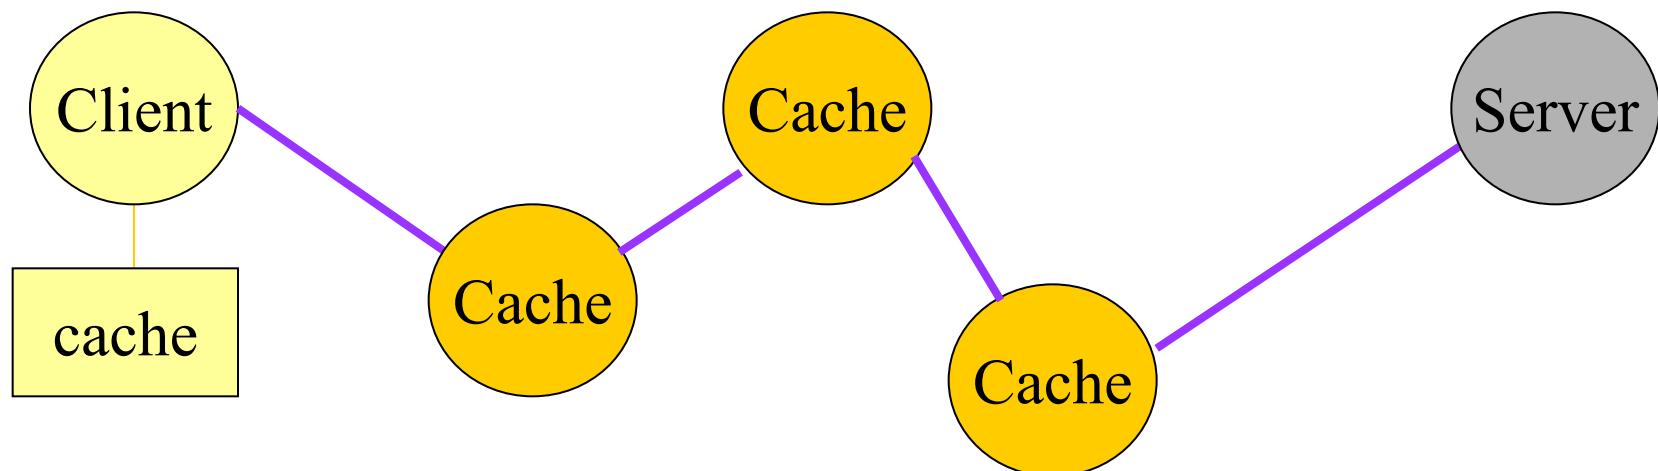
- Proxy placed close to the origin server to reduce the load on them.
Load balancing among servers of a server farm



- ## □ Interception proxy. Are invisible to the user. All messages pass through the interception proxy that can intercept them, examine the request, generate a response locally or redirect the request to another place (e.g. to the '**best replica**)

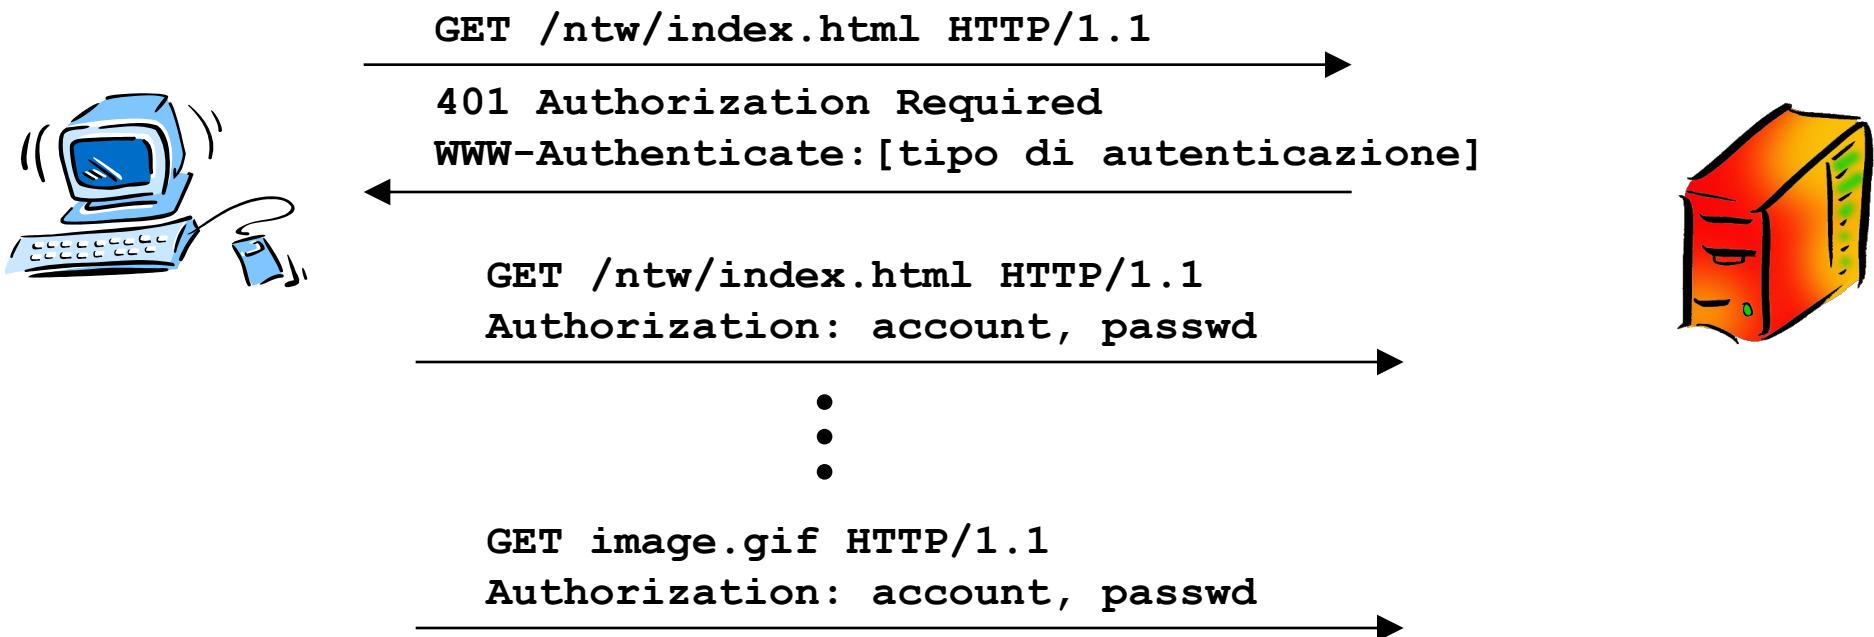
Hierarchical Caching

- The message being transported can cross several caches (e.g. at the department, ISP side, in front of the origin server -reverse proxies)
- Decreases communication costs
- Regional/national cache → high hit ratio



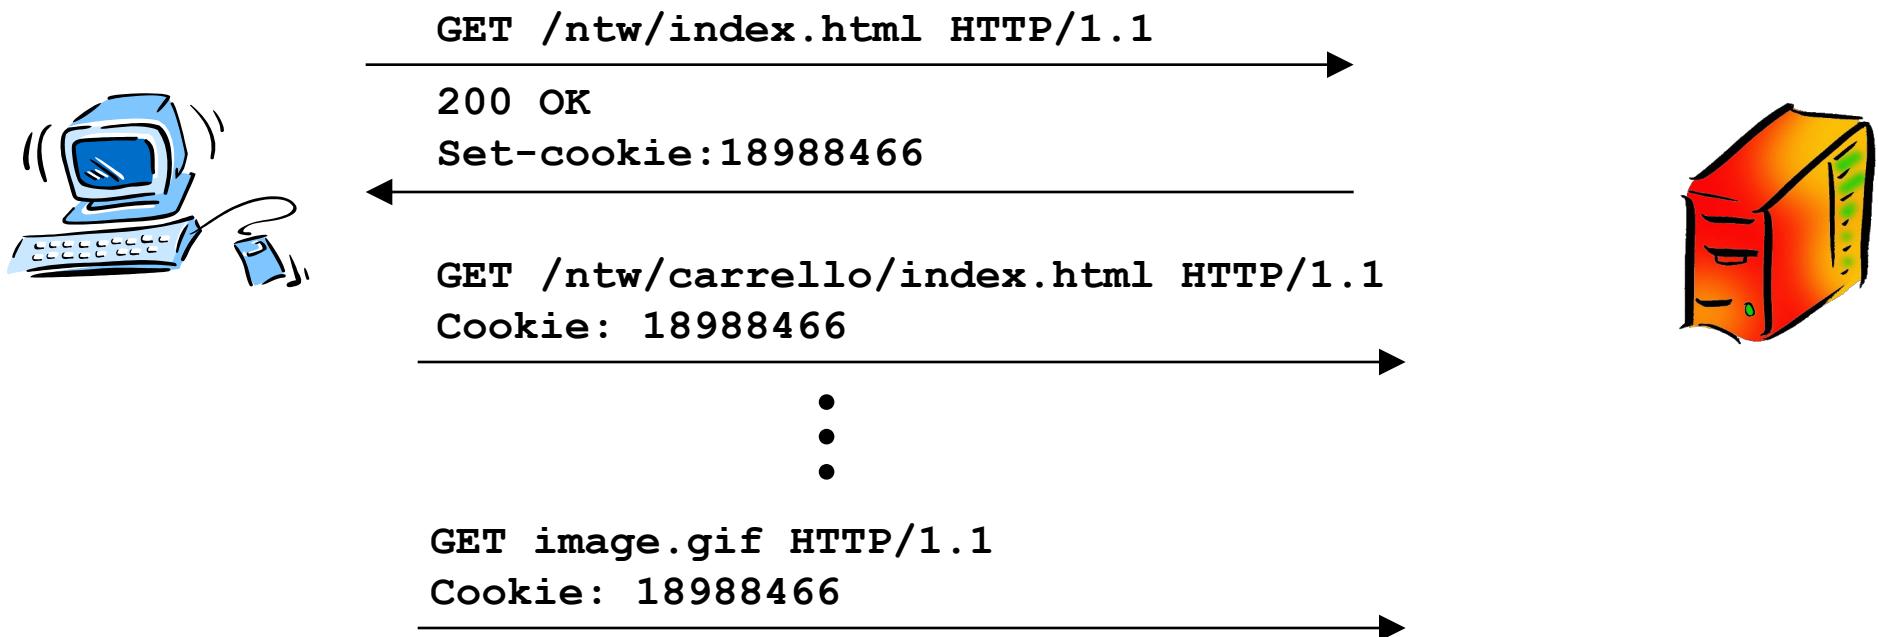
Autenticazione

- HTTP è stateless e quindi non si possono riconoscere richieste successive dello stesso utente
- in HTTP esiste un elementare meccanismo di autenticazione (account e password) che serve a riconoscere gli utenti
- Normalmente il browser memorizza passwd e account in modo da non richiedere la digitazione ogni volta



Cookie

- Esiste anche un altro modo per riconoscere richieste successive di uno stesso utente che non richiede di ricordare password
- Il numero di cookie inviato dal server viene memorizzato in un opportuno file
- Tramite i cookie si può mantenere uno stato "virtuale" per ciascun utente.

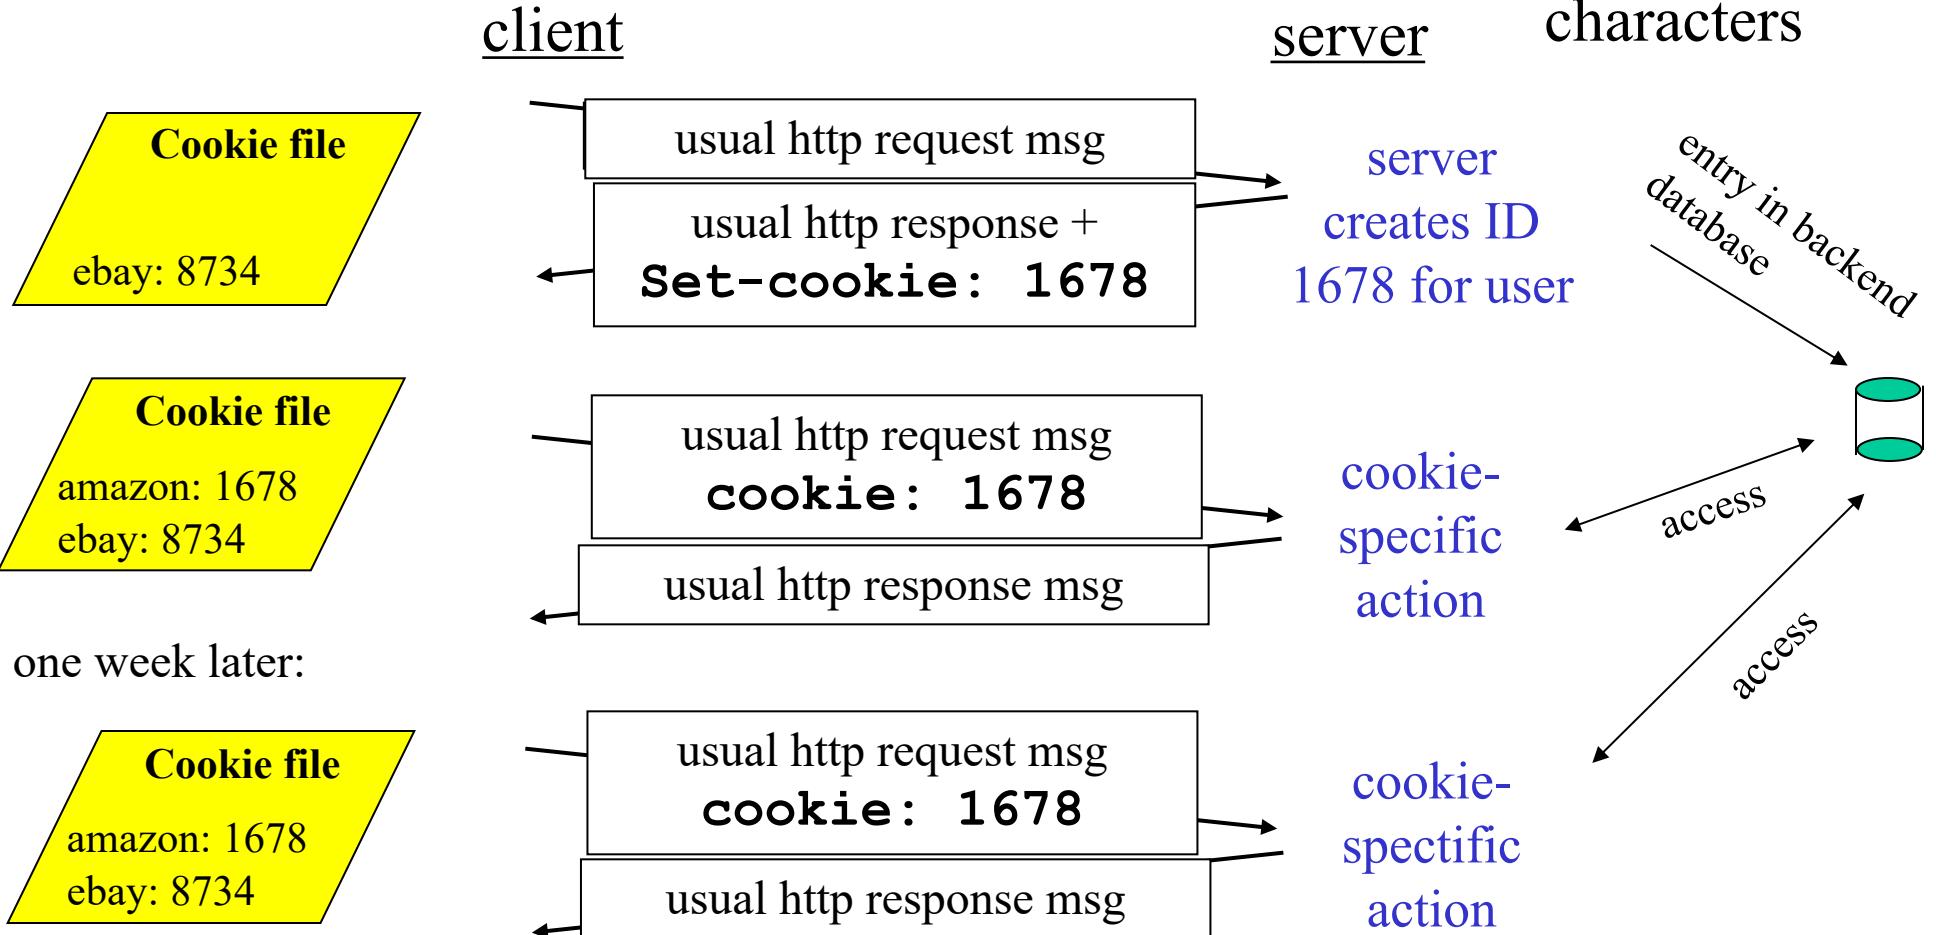


Cookie

- Permette di identificare gli utenti e mantenere delle info di stato su una transazione che richiede vari scambi di messaggi HTTP
 - Personalizzazione
 - un sito di e-commerce può personalizzare la risposta con il nome dell'utente e suggerimenti di acquisto personalizzati
 - Mantenere informazioni tra sessioni o interazioni
 - esempio: shopping cart

Cookies: keeping “state” (cont.)

Cookie=
string of
characters



Identifier to find user info on the back end database

Application Layer

Cookie

- User control on cookies
 - Whether to accept any cookies at all
 - Set a limit on the size and number of cookies
 - Limit the sites/domains from which a cookie can be accepted
 - Limit to a specific session acceptance of cookies
 - Require that cookies must originate from the same server as the current page being viewed
- Privacy concerns
 - Allow to track user behavior
 - Users typically not aware of when a cookie is sent
 - and on the use of the information he/she is providing through use of cookies
 - user profiling, shared between companies
 - Could also be sent from a Web Server different from the one user is aware to being connected to
 - e.g. due to redirection to download some objects referred to in the Main Server web page

Breve storia del Web

- 1945: Vannevar Bush in ‘As we may think’ proposes Memex to extend human memory via mechanical means. V. Bush wanted to make the existing store of knowledge (fastly growing) more accessible to mankind. ‘A memex is a device in which an individual stores all his books, records, and communications, and which is mechanized so that it maybe consulted with exceeding speed and flexibility. It is an enlarged supplement to his memory. ... Wholly new forms of encyclopedias will appear ready made with a mesh of **associative trails** running through them’
- 1965 Ted Nelson coins the term ‘hypertext’ to describe nonsequential writing that presents information as a collection of linked nodes. ‘Readers can pursue the information in a variety of ways by navigating from one node to another’
- 1970s scientific community communicate through ARPANET (exchanging e-mails, file transfers..)
- 1989 Tim Berners-Lee creates the World Wide Web (networked application that links users and services distributed across computers around the world. Users can view pages, search for information, send and receive e-mails, initiate business transactions etc. Tim Berners Lee proposed linking info presented on various machines at CERN (European Laboratory for Particle Physics, Geneva). Other systems had been created to download files (FTP)from remote computers, search for information (Gopher), WAIS (Wide Area Information Servers) etc. Why did the Web succed? 1)SIMPLE INTERFACE 2) ENHANCED FEATURES 3) USE OF HYPERTEXT

Breve storia del Web

- 1991 First browser and server introduced
- 1993 The Web consisted of around 50 servers
- 1993 First release of Mosaic browser for windows (written by Marc Andreesen and Eric Bina). The web accounted for 1% of the traffic of the Internet
- Late 1990s → Web responsible for over 75% of the Internet traffic!! Hundreds millions of users; millions of web sites. Reasons for success: graphical user interface, ease of publishing new content
- Web used for e-business, business to business, interactions between users (chatrooms and games). Web increasingly used also to access private or proprietary data.
 - Web traffic: small number of Web pages accounts for the majority of web accesses
 - from text to images, animations, video files (increased storage needed at the Web servers, increased time for downloading, increased load on the network)
 - Dynamically generated content (e.g. search engines generates pointers based on keywords)
 - Web used for banking, e-business; important personal info (e.g. credit card info transferred) → security, authentication needed

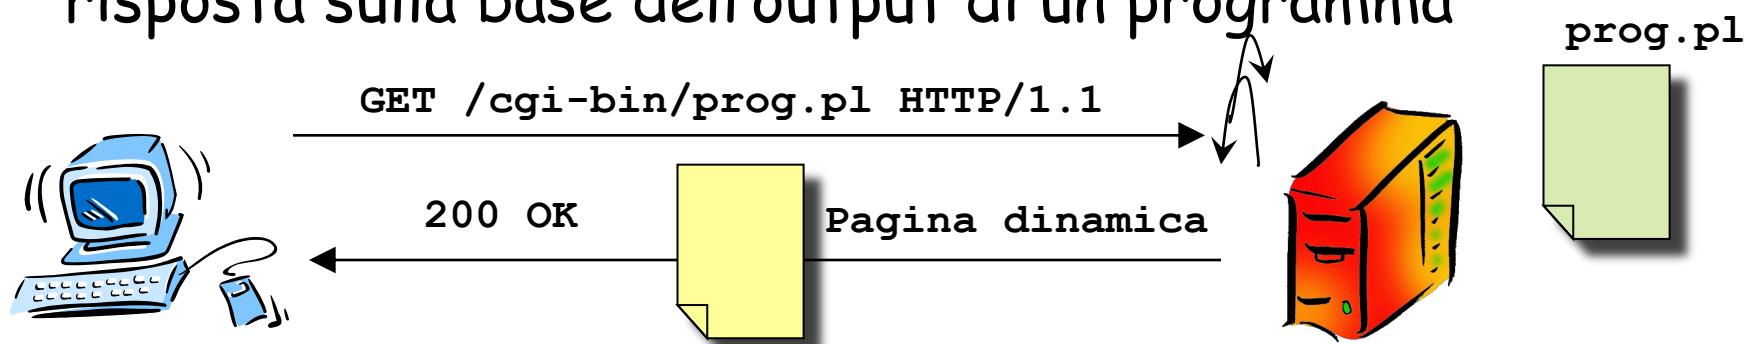
HTML

(HyperText Markup Language)

- HTTP trasferisce file e non si occupa della loro semantica
- Il funzionamento del WWW si basa sull'interpretazione di file e sulla loro visualizzazione
- Pagine di testo formattate sono trasferite come file ASCII mediante dei comandi di formattazione specificate nel linguaggio HTML
- Le pagine HTML possono contenere riferimenti ad altri oggetti che dal browser possono essere interpretate
 - come parte del documento da visualizzare (immagini)
 - Come link ad altre pagine web
- Se una pagina HTML è memorizzata nel server e viene inviata su richieste è una pagina statica

Pagine WEB dinamiche

- Se una pagina viene creata al momento della richiesta (e normalmente in base alle informazioni fornite del client) si parla di pagine dinamiche
- Se una richiesta si riferisce ad una pagina dinamica il server esamina la richiesta, esegue un programma associato a quella richiesta e genera la pagina di risposta sulla base dell'output di un programma



Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

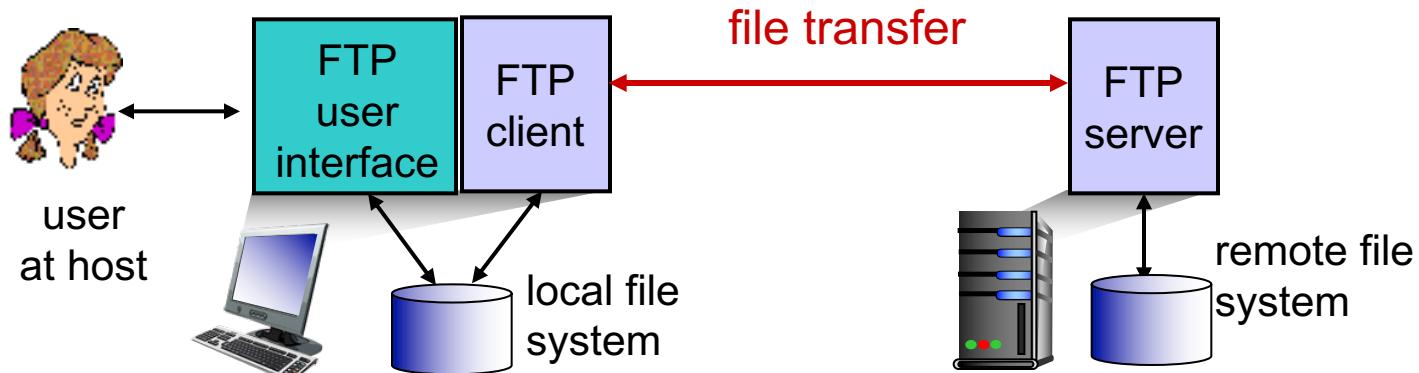
- SMTP, POP3, IMAP

2.5 DNS

2.6 P2P applications

2.7 Caching and Content Delivery Networking

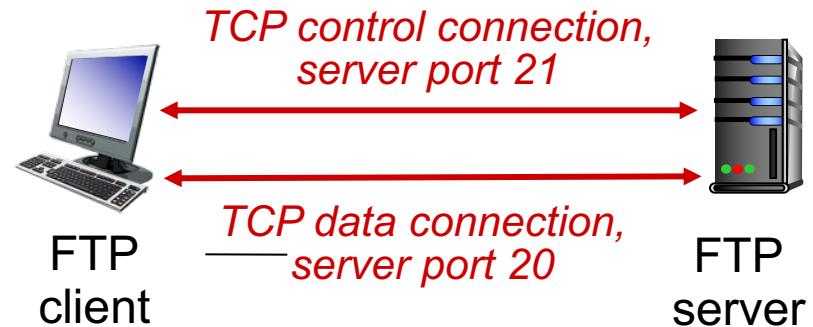
FTP: the file transfer protocol



- ❖ transfer file to/from remote host
- ❖ client/server model
 - *client*: side that initiates transfer (either to/from remote)
 - *server*: remote host
- ❖ ftp: RFC 959
- ❖ ftp server: port 21

FTP: separate control, data connections

- ❑ FTP client contacts FTP server at port 21, using TCP
- ❑ client authorized over control connection
- ❑ client browses remote directory, sends commands over control connection
- ❑ when server receives file transfer command, **server** opens 2nd TCP data connection (for file) to client
- ❑ after transferring one file, server closes data connection



- ❖ server opens another TCP data connection to transfer another file
- ❖ control connection: “*out of band*”
- ❖ FTP server maintains “state”: current directory, earlier authentication

FTP commands, responses

sample commands:

- sent as ASCII text over control channel
- USER username**
- PASS password**
- LIST** return list of file in current directory
- RETR filename** retrieves (gets) file
- STOR filename** stores (puts) file onto remote host

sample return codes

- status code and phrase (as in HTTP)
- 331 Username OK, password required**
- 125 data connection already open; transfer starting**
- 425 Can't open data connection**
- 452 Error writing file**

Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

- SMTP, POP3, IMAP

2.5 DNS

2.6 P2P applications

2.7 Caching and Content Delivery Networking

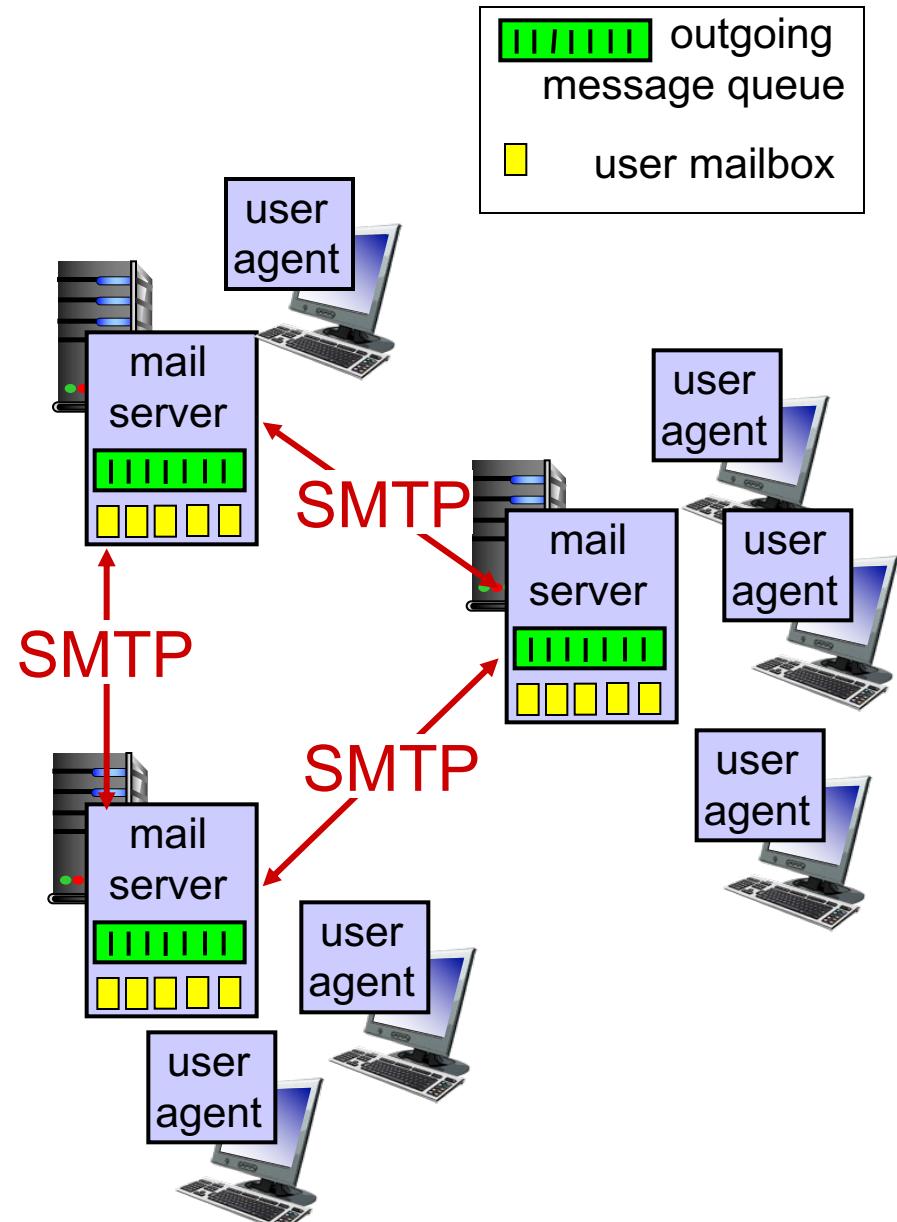
Electronic mail

Three major components:

- user agents
- mail servers
- simple mail transfer protocol: SMTP

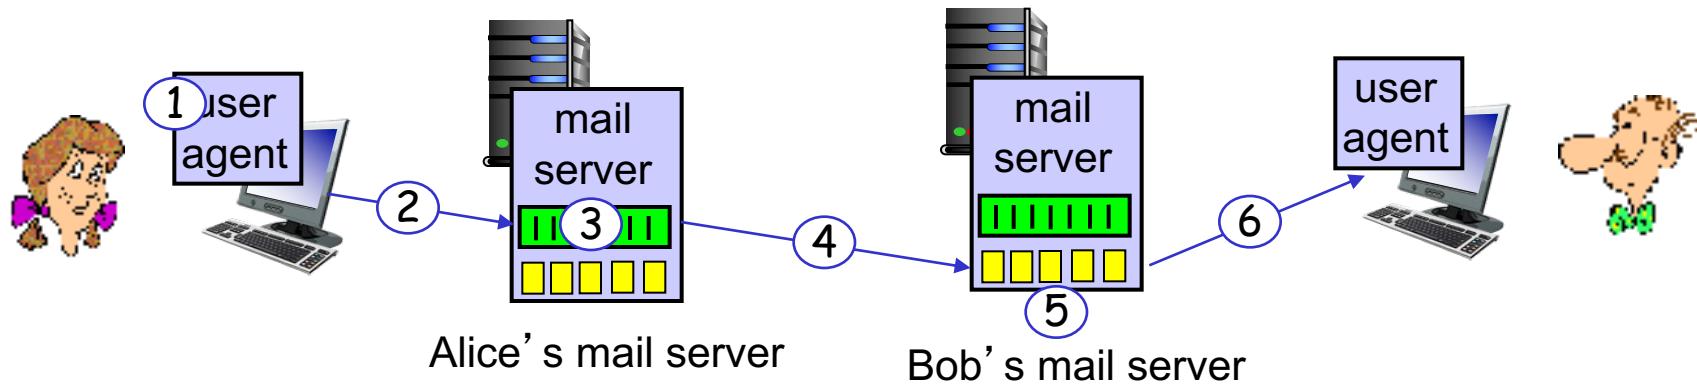
User Agent

- a.k.a. “mail reader”
- composing, editing, reading mail messages
- e.g., Outlook, Thunderbird, iPhone mail client
- outgoing, incoming messages stored on server



Scenario: Alice sends message to Bob

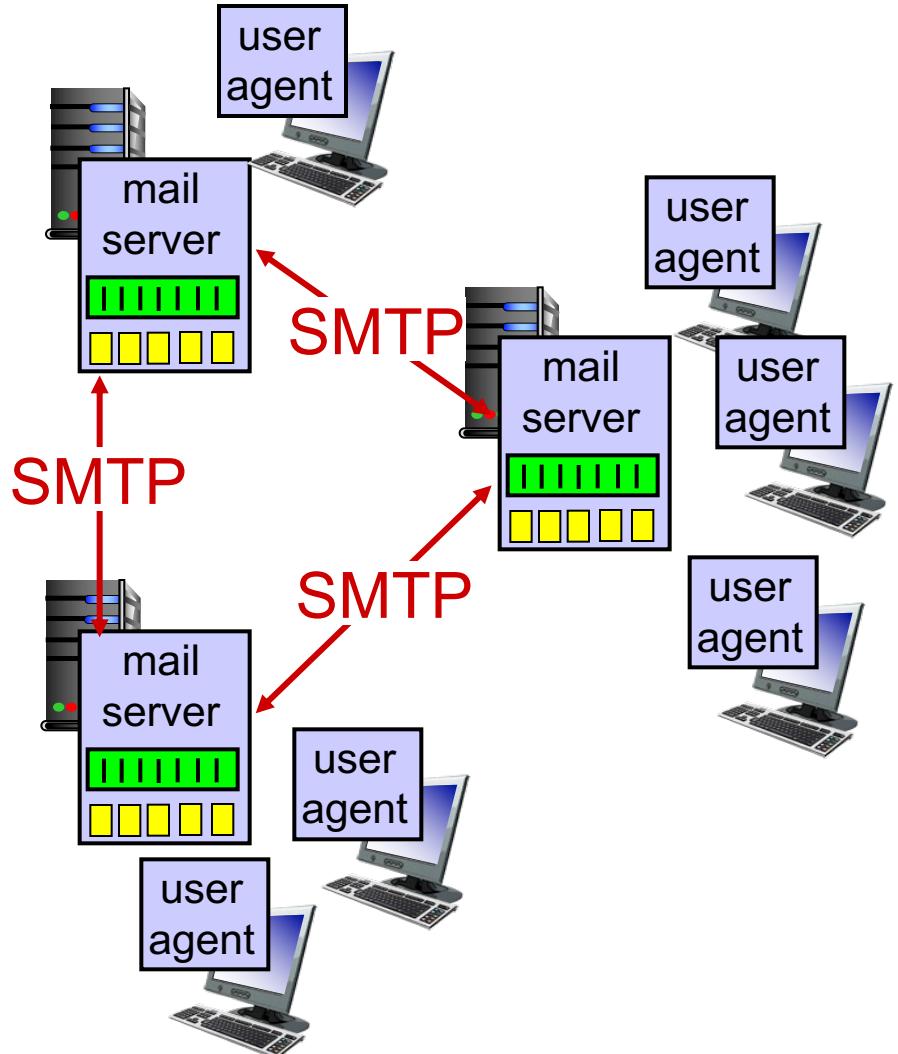
- 1) Alice uses UA to compose message “to” bob@someschool.edu
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



Electronic mail: mail servers

mail servers:

- ❑ *mailbox* contains incoming messages for user
- ❑ *message queue* of outgoing (to be sent) mail messages
- ❑ *SMTP protocol* between mail servers to send email messages
 - client: sending mail server
 - “server”: receiving mail server



Electronic Mail: SMTP [RFC 2821]

- ❑ uses TCP to reliably transfer email message from client to server, port 25
- ❑ direct transfer: sending server to receiving server
- ❑ three phases of transfer
 - handshaking (greeting)
 - transfer of messages
 - closure
- ❑ command/response interaction (like HTTP, FTP)
 - commands: ASCII text
 - response: status code and phrase
- ❑ messages must be in 7-bit ASCII

Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

SMTP: final words

- SMTP uses persistent connections
- SMTP requires message (header & body) to be in 7-bit ASCII
- SMTP server uses CRLF.CRLF to determine end of message

comparison with HTTP:

- HTTP: pull
- SMTP: push
- both have ASCII command/response interaction, status codes
- HTTP: each object encapsulated in its own response msg
- SMTP: multiple objects sent in multipart msg

Mail message format

SMTP: protocol for
exchanging email msgs

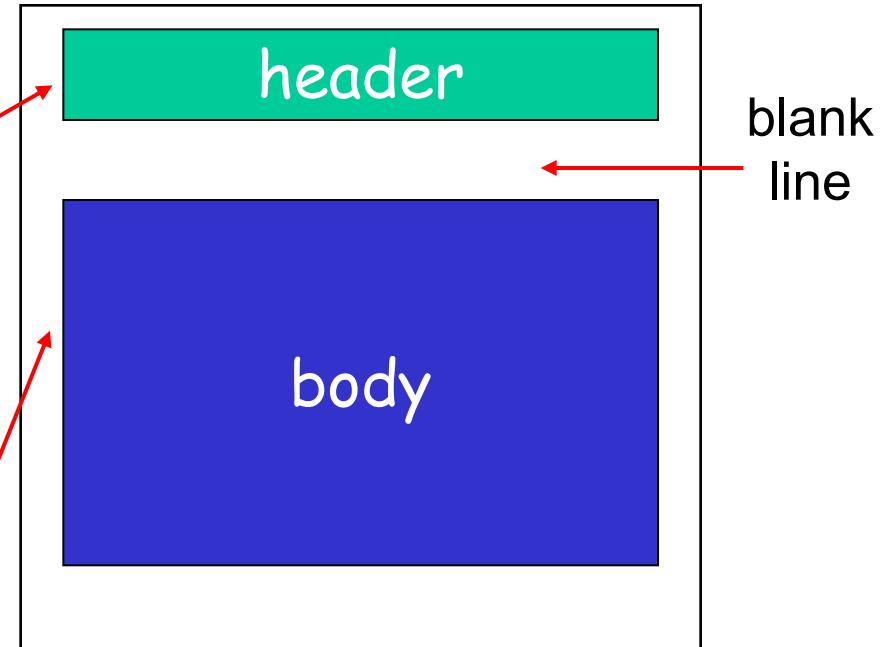
RFC 822: standard for text
message format:

- ❑ header lines, e.g.,

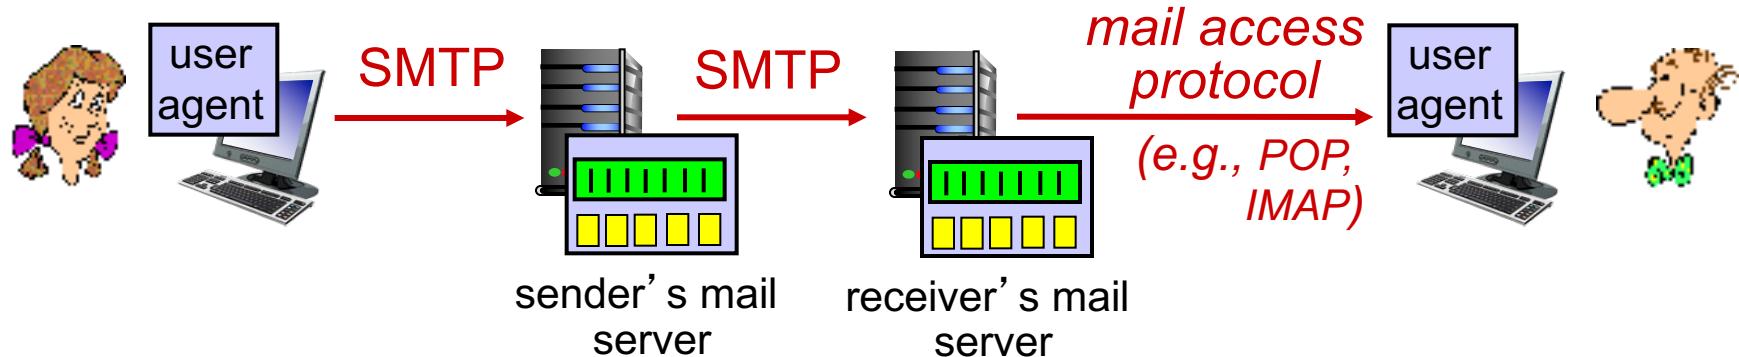
- To:
 - From:
 - Subject:

different from SMTP MAIL
FROM, RCPT TO:
commands!

- ❑ Body: the “message”
 - ASCII characters only



Mail access protocols



- **SMTP:** delivery/storage to receiver's server
- mail access protocol: retrieval from server
 - **POP:** Post Office Protocol [RFC 1939]: authorization, download
 - **IMAP:** Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored msgs on server
 - **HTTP:** gmail, Hotmail, Yahoo! Mail, etc.

POP3 protocol

authorization phase

- client commands:
 - **user**: declare username
 - **pass**: password
- server responses
 - **+OK**
 - **-ERR**

transaction phase, client:

- list**: list message numbers
- retr**: retrieve message by number
- dele**: delete
- quit**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

POP3 (more) and IMAP

more about POP3

- ❑ previous example uses POP3 “download and delete” mode
 - Bob cannot re-read email if he changes client
- ❑ POP3 “download-and-keep”: copies of messages on different clients
- ❑ POP3 is stateless across sessions

IMAP

- ❑ keeps all messages in one place: at server
- ❑ allows user to organize messages in folders
- ❑ keeps user state across sessions:
 - names of folders and mappings between message IDs and folder name

Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 FTP

2.4 electronic mail

- SMTP, POP3, IMAP

2.5 DNS

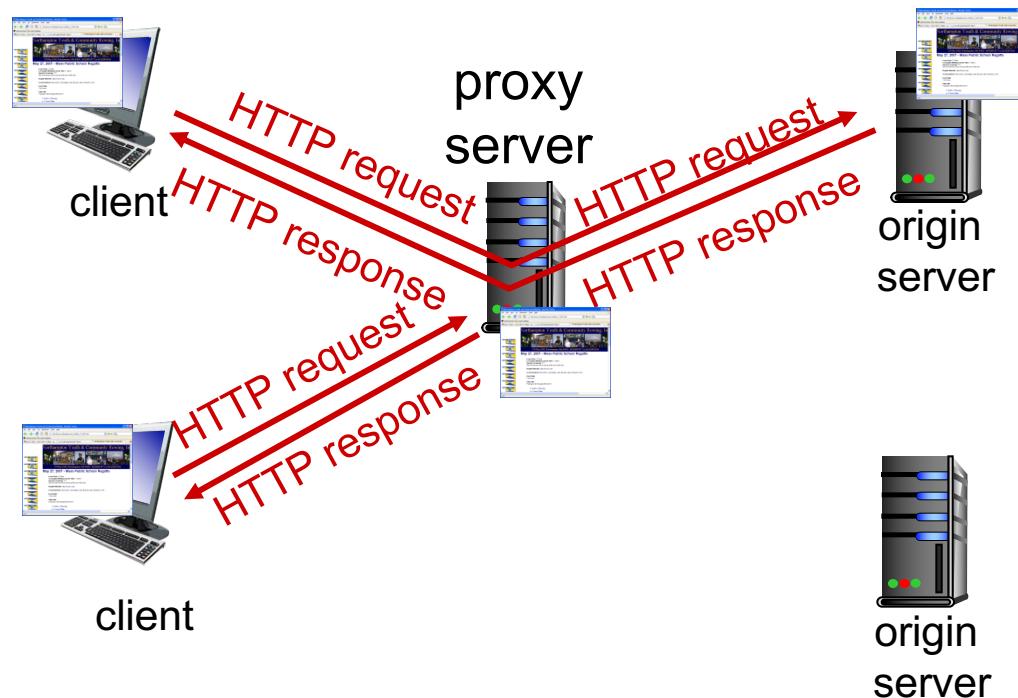
2.6 P2P applications

2.7 Caching and Content Delivery Networking

Web caches (proxy server)

goal: satisfy client request without involving origin server

- ❑ user sets browser: Web accesses via cache
- ❑ browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



More about Web caching

- cache acts as both client and server
 - server for original requesting client
 - client to origin server
- typically cache is installed by ISP (university, company, residential ISP)

why Web caching?

- reduce response time for client request
- reduce traffic on an institution's access link
- Internet dense with caches: enables “poor” content providers to effectively deliver content (so too does P2P file sharing)

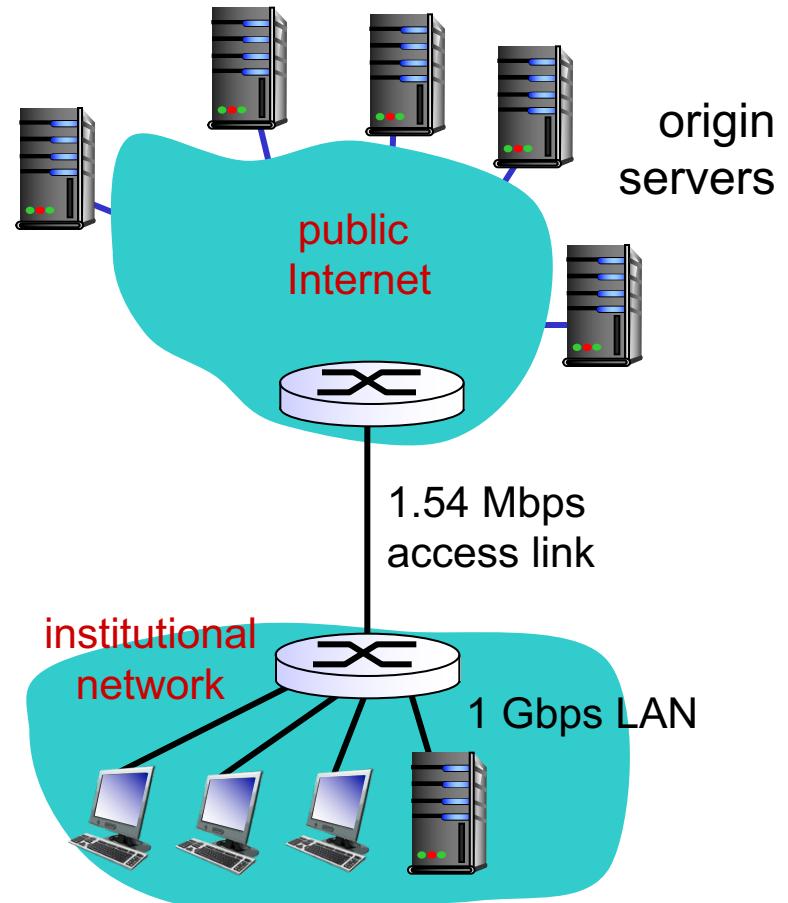
Caching example:

assumptions:

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

consequences: *problem!*

- LAN utilization: 0,15%
- access link utilization = *99%*
- total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + usecs



Caching example: fatter access link

assumptions:

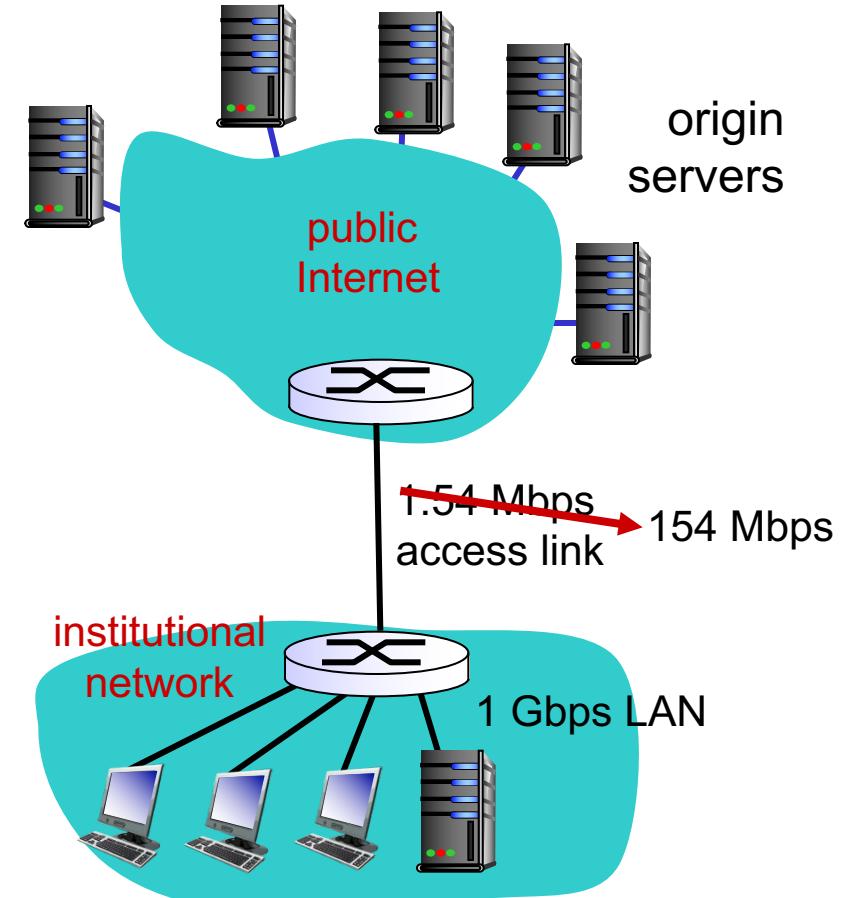
- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

154 Mbps

9.9%

consequences:

- LAN utilization: 0.15%
- access link utilization = 9.9%
- total delay = Internet delay + access delay + LAN delay
 $= 2 \text{ sec} + \cancel{\text{minutes}} + \cancel{\text{usecs}}$ → msecs



Cost: increased access link speed (not cheap!)

Caching example: install local cache

assumptions:

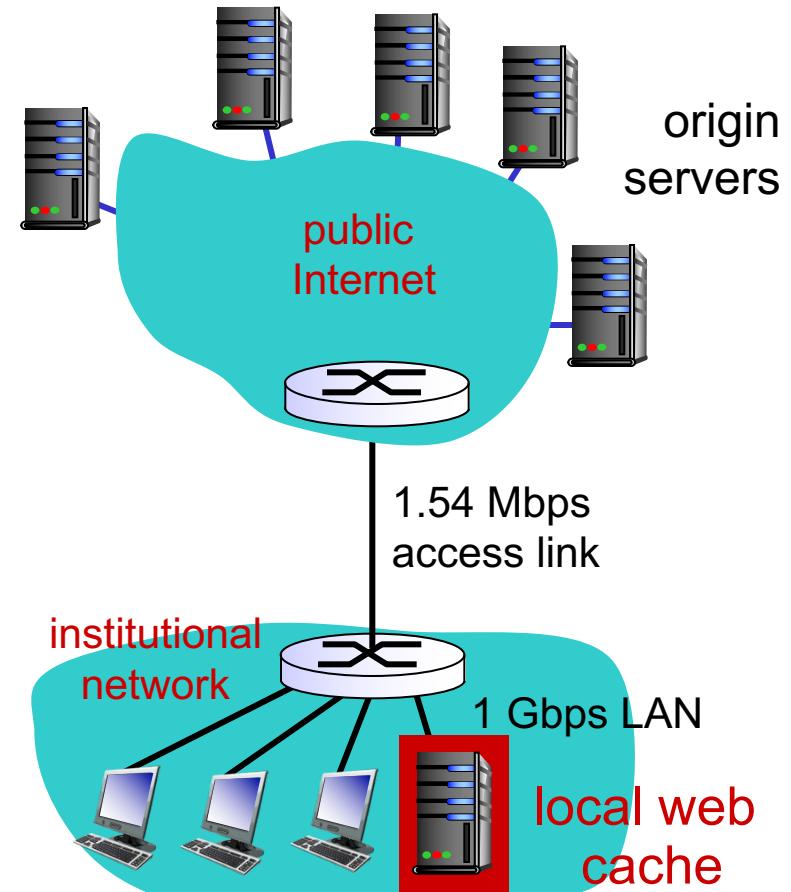
- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

consequences:

- LAN utilization: 0,15%
- access link utilization = 100%
- total delay = Internet delay + access delay + LAN delay
= 2 sec + n · ?

How to compute link utilization, delay?

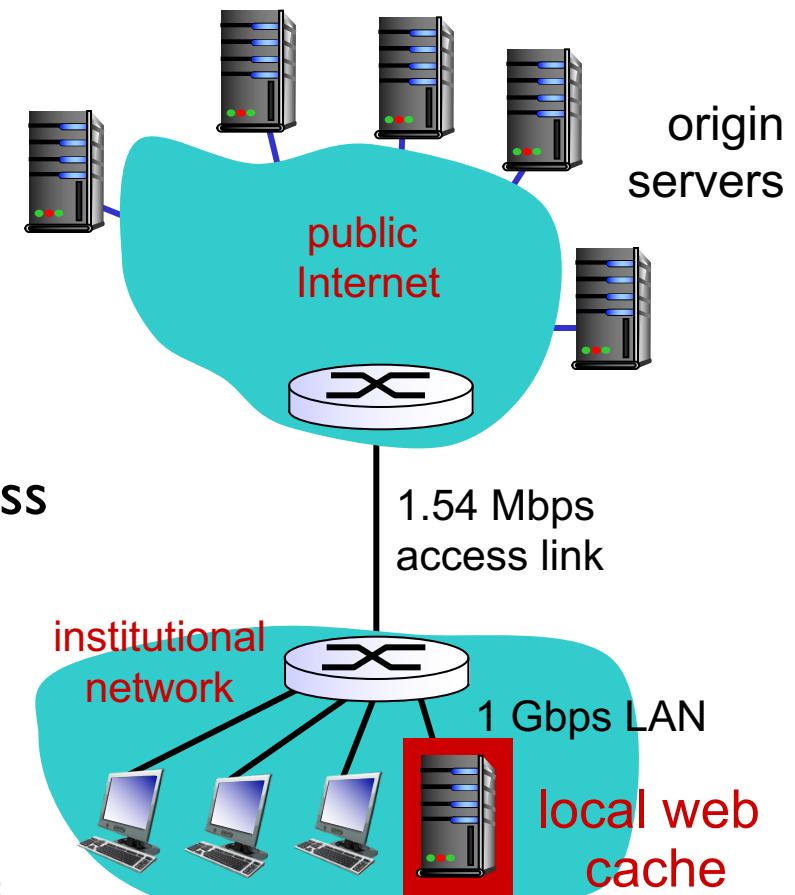
Cost: web cache (cheap!)



Caching example: install local cache

Calculating access link utilization, delay with cache:

- ❑ suppose cache hit rate is 0.4
 - 40% requests satisfied at cache,
60% requests satisfied at origin
- access link utilization:
 - 60% of requests use access link
- data rate to browsers over access link
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
 - utilization = $0.9 / 1.54 = .58$
- total delay
 - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
 - $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
 - less than with 154 Mbps link (and cheaper too!)



DNS: domain name system

people: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

- IP address (32 bit in IPv4) - used for addressing datagrams
- “name”, e.g., www.yahoo.com - used by humans

Q: how to map between IP address and name, and vice versa ?

Domain Name System:

- *distributed database* implemented in hierarchy of many *name servers*
- *application-layer protocol:* hosts, name servers communicate to *resolve* names (address/name translation)
 - note: core Internet function, implemented as application-layer protocol
 - complexity at network’s “edge”

DNS: services, structure

DNS services

- hostname to IP address translation
- host aliasing
 - canonical, alias names
- mail server aliasing
- load distribution
 - replicated Web servers: many IP addresses correspond to one name

why not centralize DNS?

DNS: services, structure

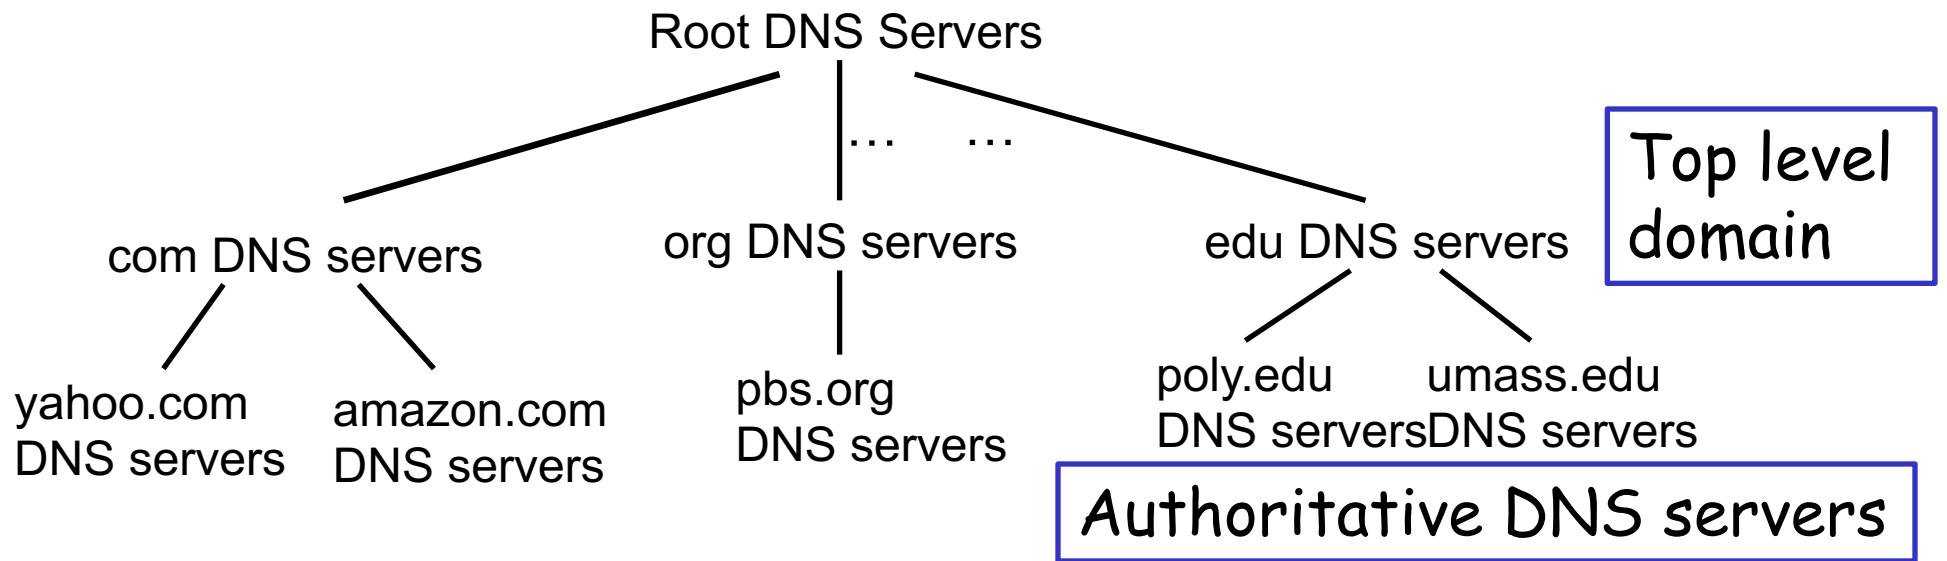
DNS services

- hostname to IP address translation
- host aliasing
 - canonical, alias names
- mail server aliasing
- load distribution
 - replicated Web servers: many IP addresses correspond to one name

why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

DNS: a distributed, hierarchical database

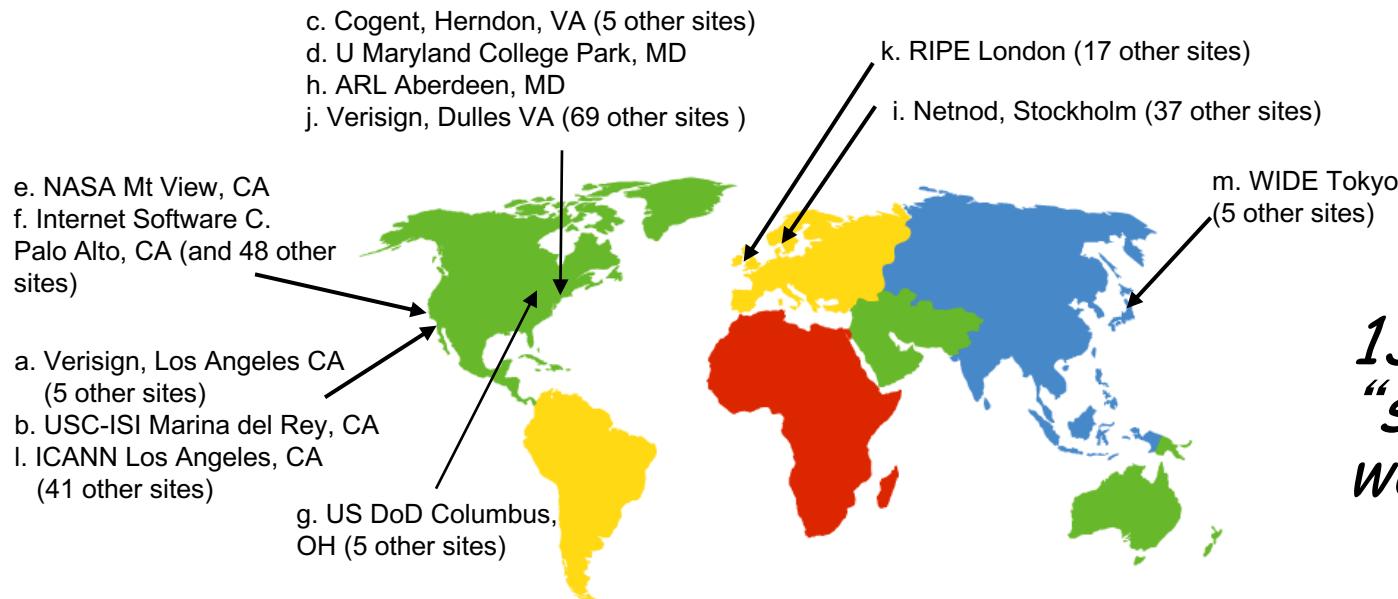


client wants IP for www.amazon.com; 1st approx:

- client queries root server to find com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

DNS: root name servers

- contacted by local name server that can not resolve name
- root name server:
 - could contacts authoritative name server if name mapping not known (in recursive queries)
 - gets mapping
 - returns mapping to local name server



*13 root name
“servers”
worldwide*

TLD, authoritative servers

top-level domain (TLD) servers:

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp, eu
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD

authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

Local DNS name server

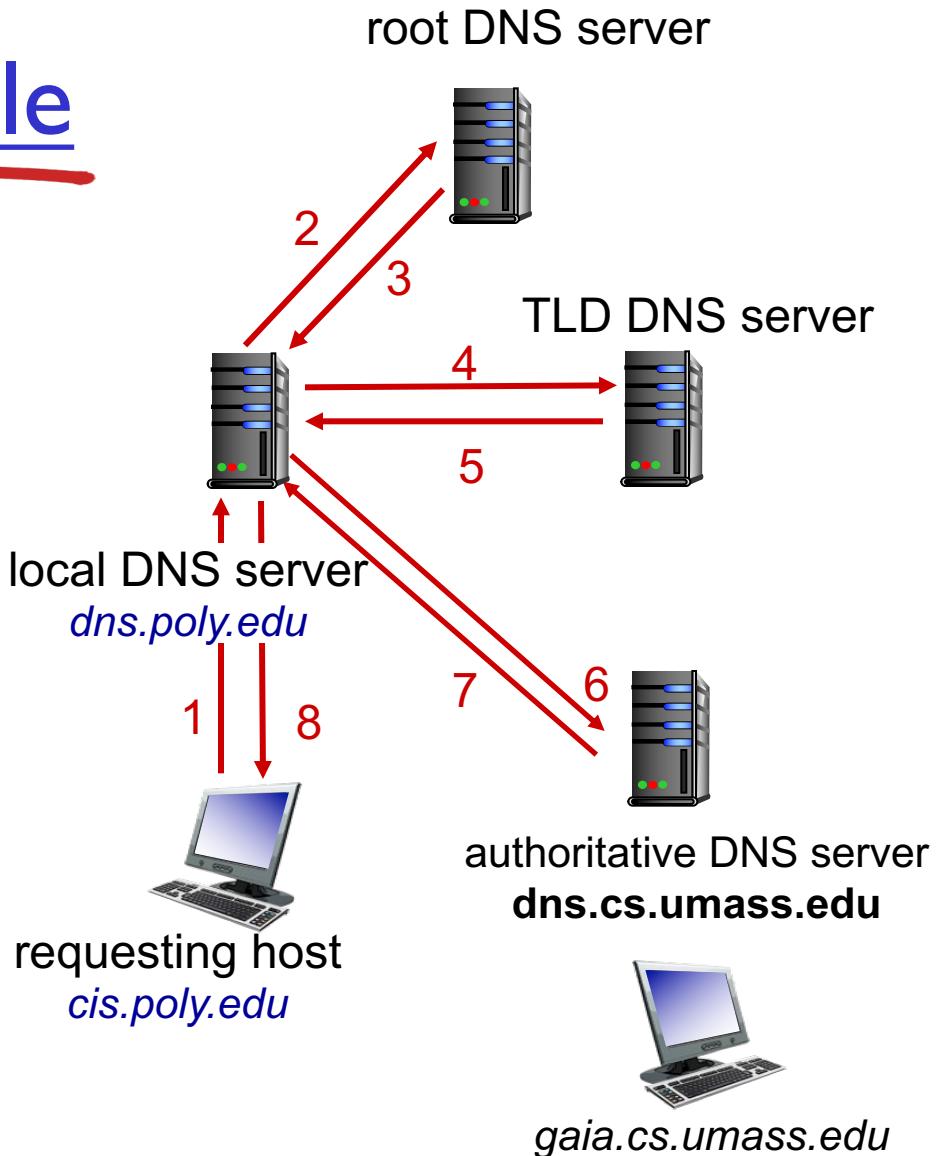
- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
 - also called “default name server”
- when host makes DNS query, query is sent to its local DNS server
 - has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy, forwards query into hierarchy

DNS name resolution example

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

iterated query:

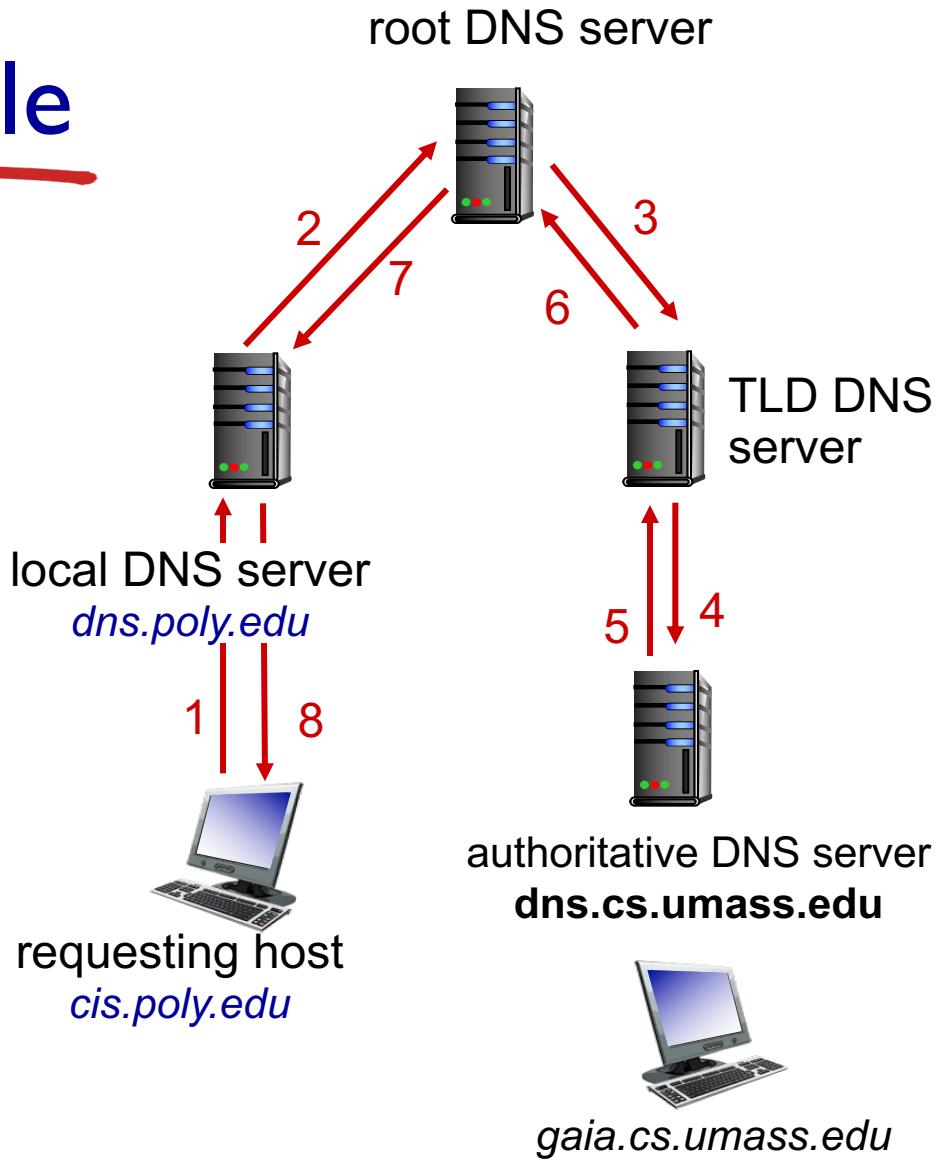
- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



DNS name resolution example

recursive query:

- ❖ puts burden of name resolution on contacted name server
- ❖ heavy load at upper levels of hierarchy?



DNS: caching, updating records

- once (any) name server learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time (TTL)
 - TLD servers typically cached in local name servers
 - thus root name servers not often visited
- cached entries may be *out-of-date* (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- update/notify mechanisms proposed IETF standard
 - RFC 2136

DNS records

DNS: distributed db storing resource records (RR)

RR format: (**name**, **value**, **type**, **ttl**)

type=A

- **name** is hostname
- **value** is IP address
(relay.bar.foo.com, 145.37.93.126,A)

type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain
(foo.com,dns.foo.com,NS)

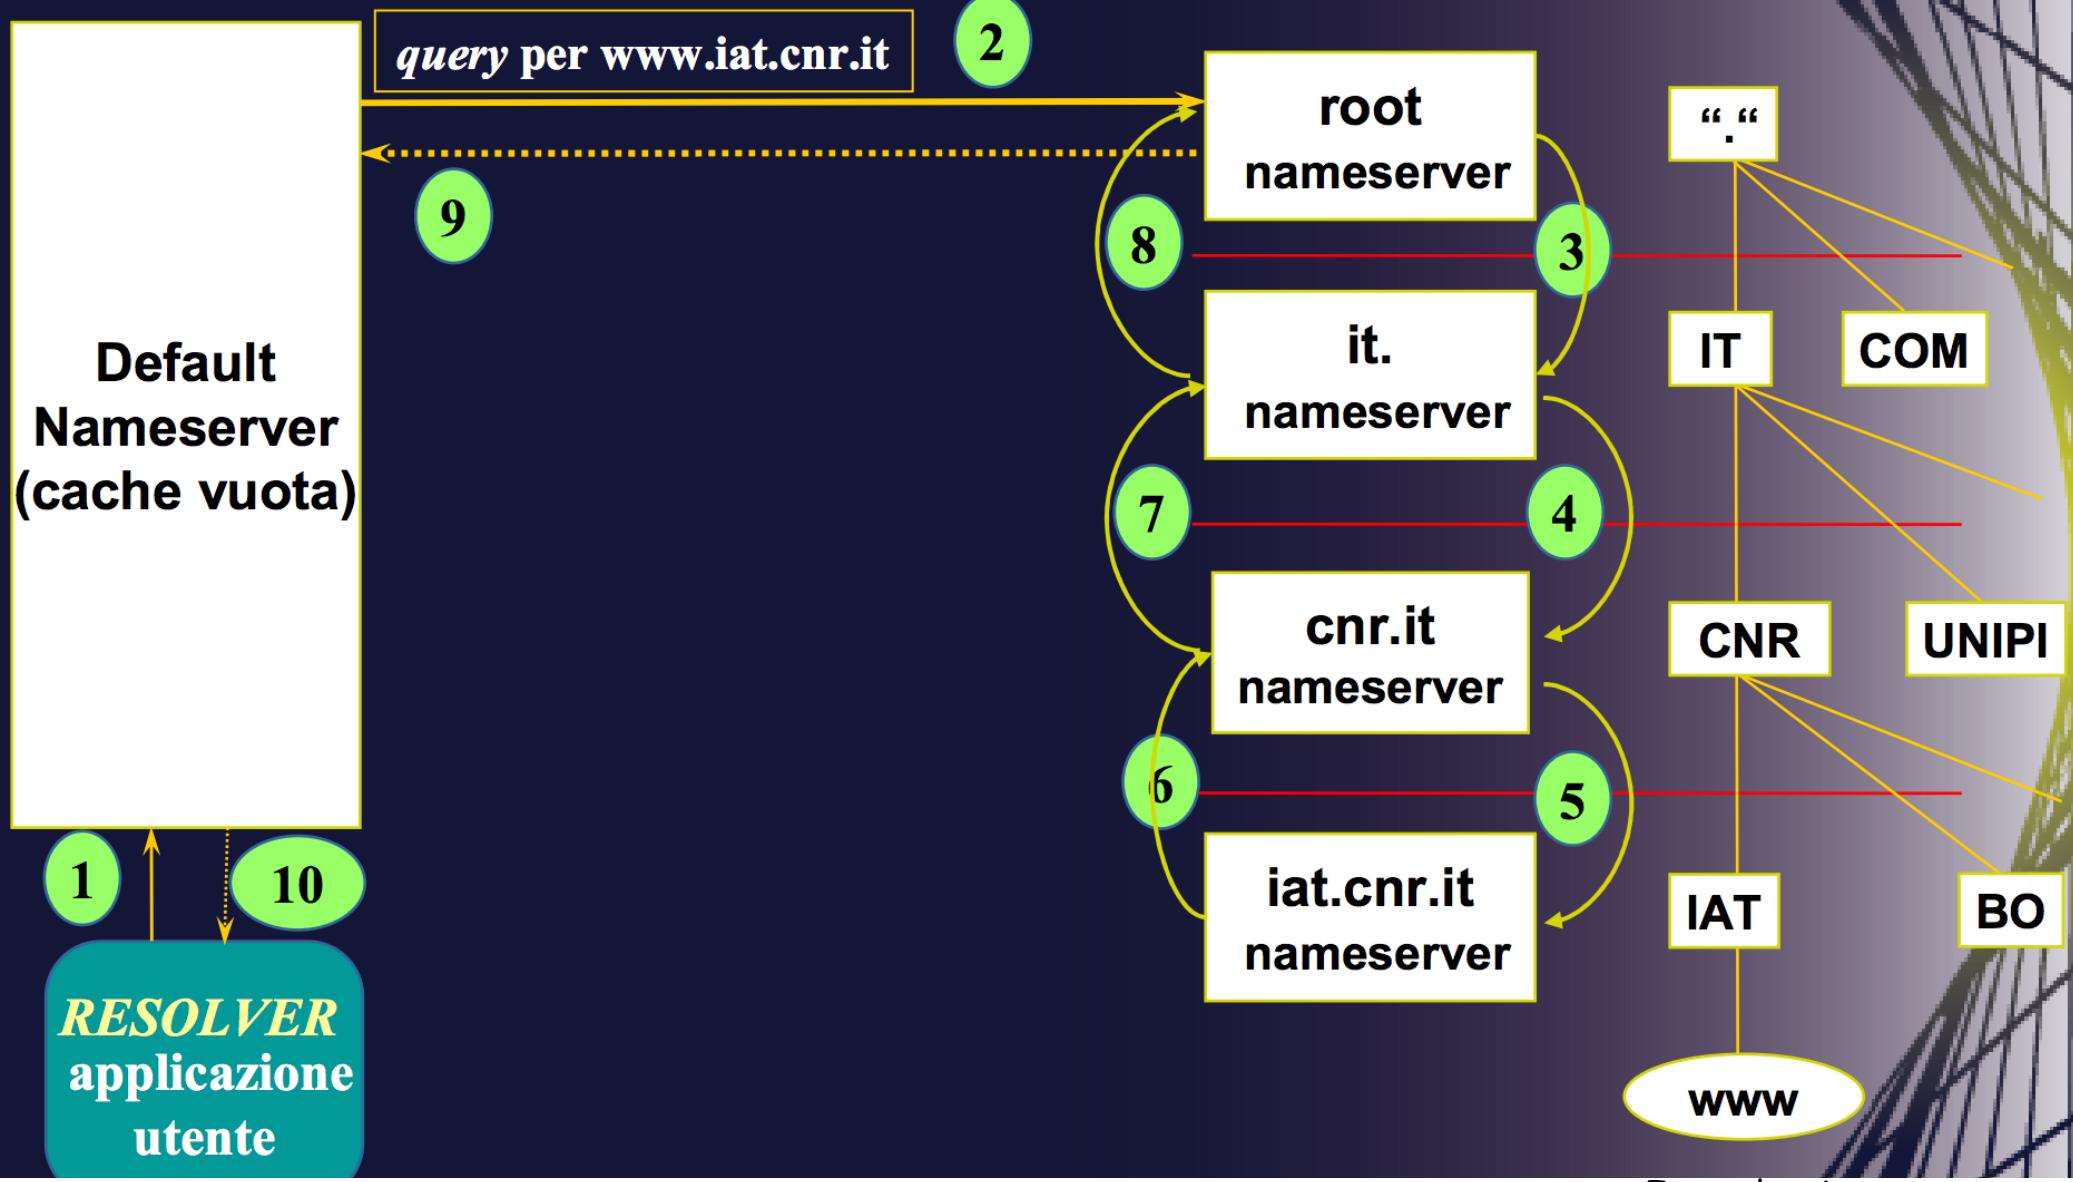
type=CNAME

- **name** is alias name for some “canonical” (the real) name
- www.ibm.com is really servereast.backup2.ibm.com
- **value** is canonical name

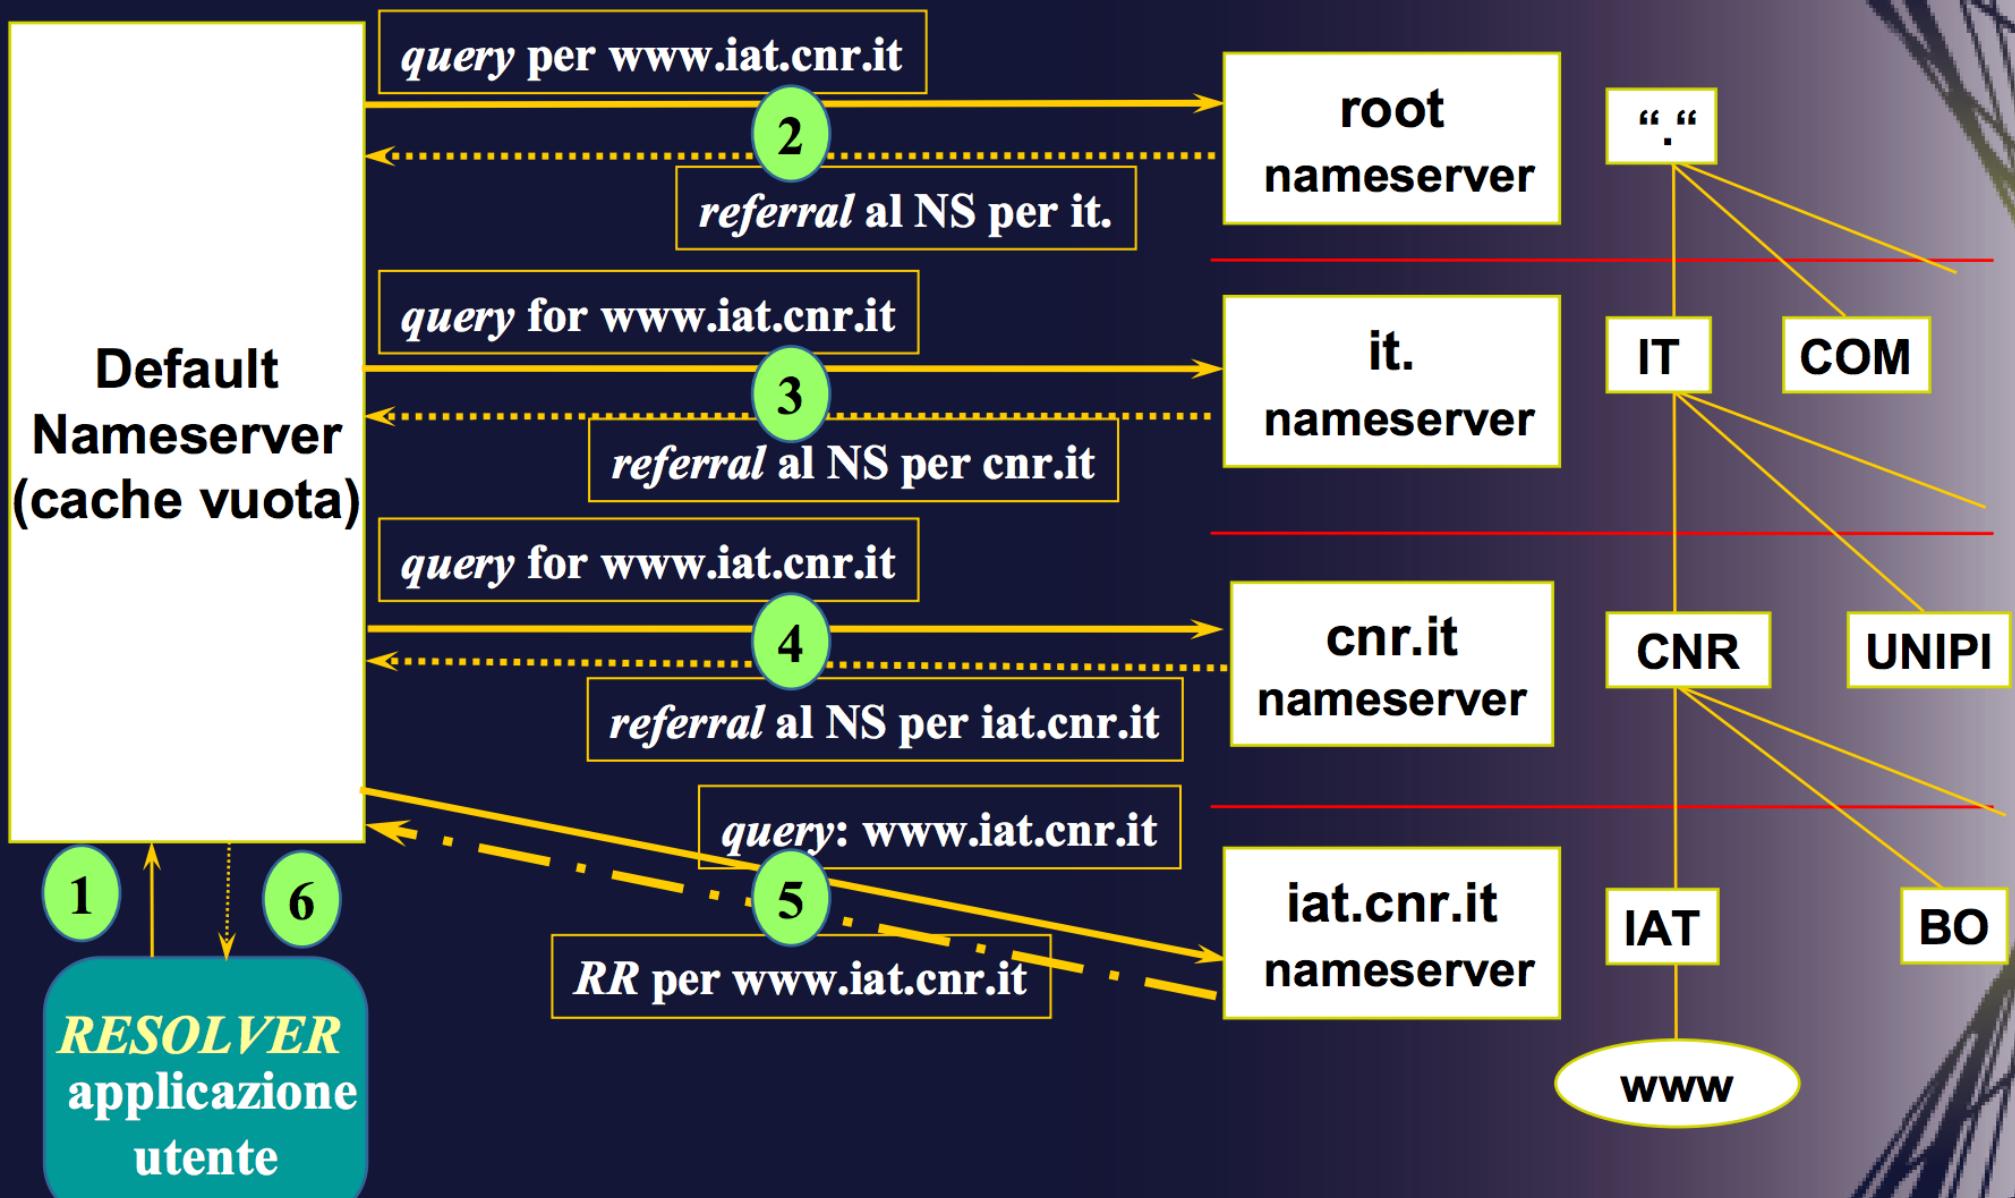
type=MX

- **value** is name of mailserver associated with **name**

Il processo ricorsivo di risoluzione dei nomi



Il processo iterativo di risoluzione dei nomi

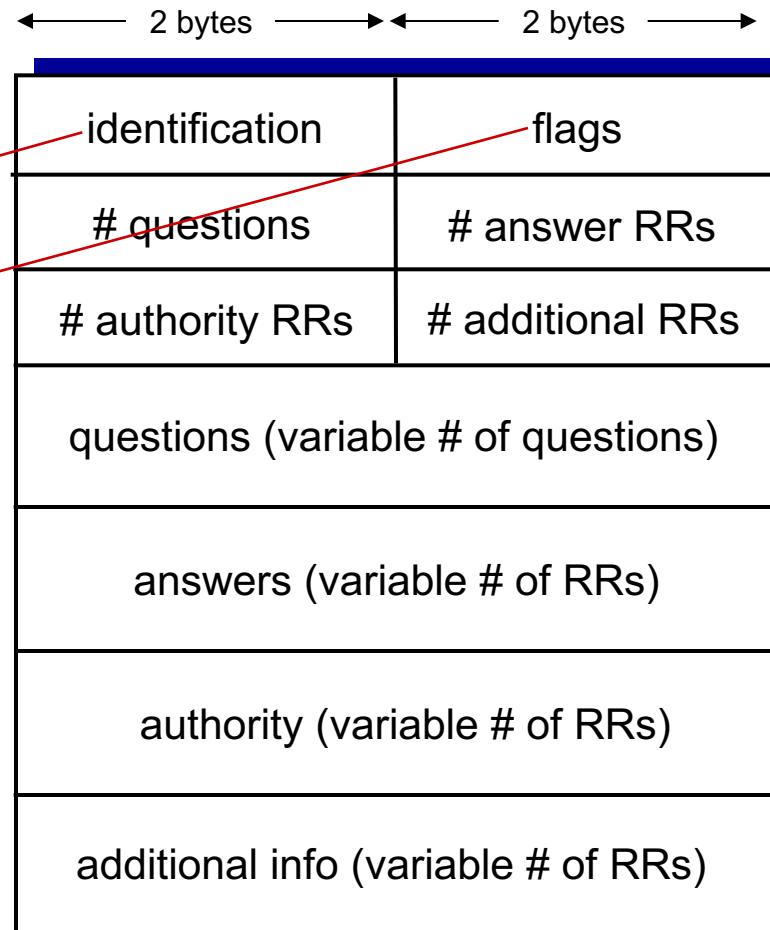


DNS protocol, messages

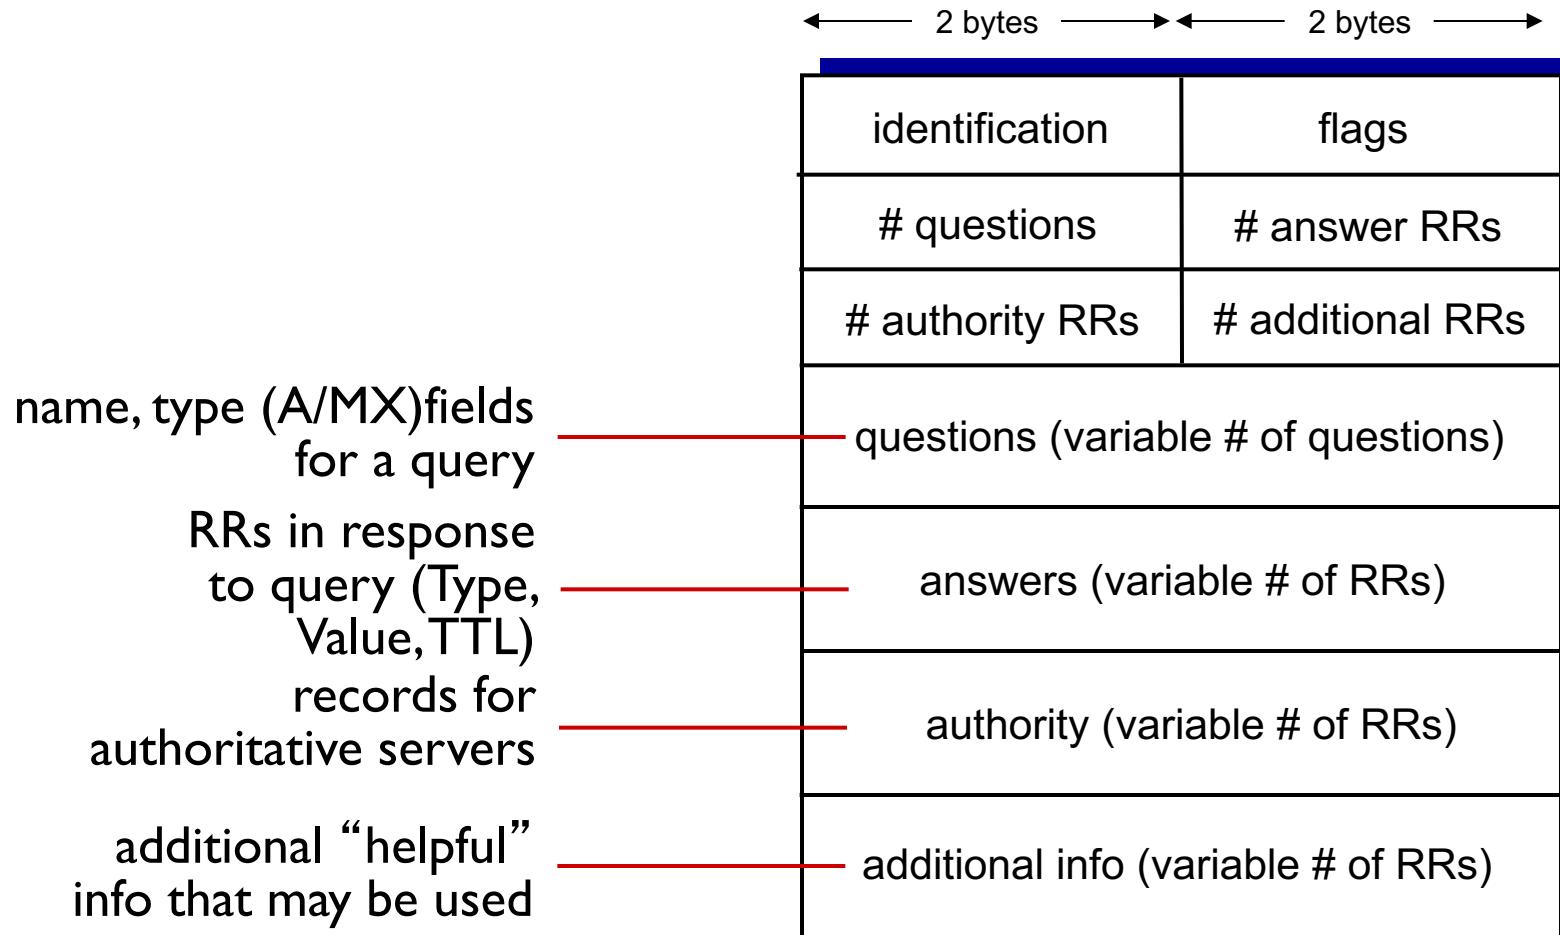
- ❑ *query* and *reply* messages, both with same *message format*

msg header

- ❖ **identification:** 16 bit #
for query, reply to query
uses same #
- ❖ **flags:**
 - query or reply
 - recursion desired
 - recursion available
 - reply is authoritative



DNS protocol, messages



Inserting records into DNS

- ❑ example: new startup “Network Utopia”
- ❑ register name networkuptopia.com at *DNS registrar* (e.g., Network Solutions)
 - provide names, IP addresses of authoritative name server (primary and secondary)
 - registrar inserts two RRs into .com TLD server:
(`networkutopia.com`, `dns1.networkutopia.com`, NS)
(`dns1.networkutopia.com`, `212.212.212.1`, A)
- ❑ create authoritative server type A record for `www.networkuptopia.com`; type MX record for `networkutopia.com`

Perche' UDP?

- Less overhead
 - Messaggi corti
 - Tempo per set-up connessione di TCP lungo
 - Un unico messaggio deve essere scambiato tra una coppia di server (nella risoluzione contattati diversi server—se si usasse TCP ogni volta dovremmo mettere su la connessione!!)
- Se un messaggio non ha risposta entro un timeout?
 - Semplicemente viene riinviauto dal resolver (problema
 - Risolto dallo strato applicativo)

Porta usata per il DNS: 53!!